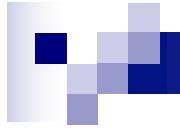


CS412

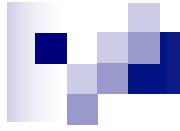
Machine Learning

Introduction – 1b

Berrin Yanikoglu
Sabancı Üniversitesi

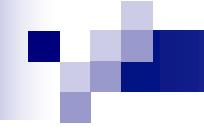


Different Learning Paradigms



A Loose Taxonomy of ML

- **Supervised learning**—where the algorithm learns a model that maps a set of inputs to a set of desired outputs, using **labelled data**.
- **Unsupervised learning**—where the algorithm discovers structure, similar patterns, sub-spaces, etc.
- **Semi-supervised or weakly supervised learning**—an intermediate supervised and unsupervised setting
- **Reinforcement learning**—where the algorithm learns a policy or a model in order to maximize a reward value without explicit feedback at intermediate steps.



Supervised Learning

- Training set of (\mathbf{x}, y) pairs where y is the desired target label given by “teacher”
$$D = \{ (\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_N, y_N) \}$$
- Learn the unknown mapping $y=f(\mathbf{x})$ so as to predict the target value for a future/unseen input \mathbf{x} .

Classification: Label y is a category

- Gender classification ($y \in \{0,1\}$)
- Spam classification ($y \in \{0,1\}$)
- Digit classification ($y \in \{0..9\}$)

Regression: Label y is a scalar ($y \in \mathbb{R}$)

- Credit scoring
- Sales prediction



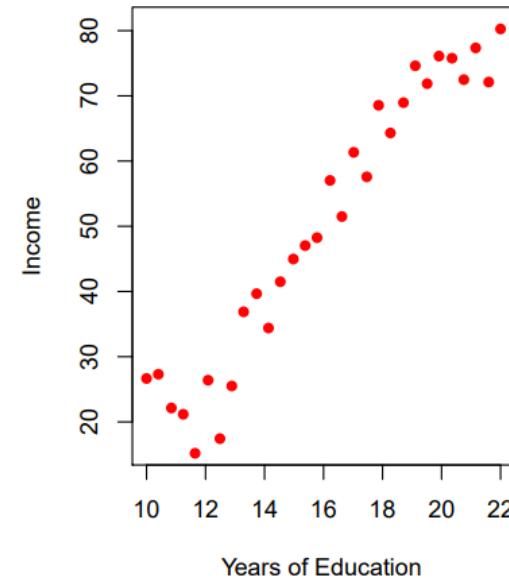
Classification

- Target to learn is a **class label** (binary or categorical label)
- Examples:
 - Medical diagnosis: mammogram **cancerous or not?**
 - Credit card applications or transactions: **fraud or not?**
 - Spam filtering in email: **spam or not?**
 - Mail categorization: **inbox or junk or promotions of forum?**
 - Speech understanding: **which word is spoken?**
 - Handwritten letters: **which word/letter is this?**

Regression

- Target to learn is a **scalar**
- Examples:
 - **Credit scoring**: given the bank, employment... data about a person, what should be their credit score
 - **House price prediction**: Given features of a house, what is a fair price?
 - **Sentiment Analysis**: Given a text, what is the scalar value of the sentiment in that text

How does income relate to the number of years of education?



MNIST Dataset

Imagine the **digit image classification** problem

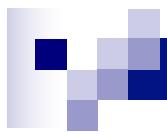
- historically called [Optical Character Recognition \(OCR\)](#)

We want to **build (learn)** a machine learning **model** that can take a binary image and output its digit (**classify** it)

We will be using this very well-known dataset in our first homeworks.



Figure 1: MNIST Dataset



MNIST Dataset

Imagine the digit image classification problem

- historically called Optical Character Recognition (OCR)

We want to **build (learn)** a machine learning **model** that can take a binary image and output its digit (**classify** it)

We will be using this very well-known dataset in our first homeworks.

0	8	7	6	4	6	9	7	2	1	5	1	4	6	
0	1	2	3	4	4	6	2	9	3	0	1	2	3	4
0	1	2	3	4	5	6	7	0	1	2	3	4	5	0
7	4	2	0	9	1	2	8	9	1	4	0	9	5	0
0	2	7	8	4	8	0	7	7	1	1	2	9	3	6
5	3	9	4	2	7	2	3	8	1	2	9	8	8	7
2	9	1	6	0	1	7	1	1	0	3	4	2	6	4
7	7	6	3	6	7	4	2	7	4	9	1	0	6	8
2	4	1	8	3	5	5	5	3	5	9	7	4	8	5

Figure 1: MNIST Dataset

Black pixels are 0,
White pixels are 1

or

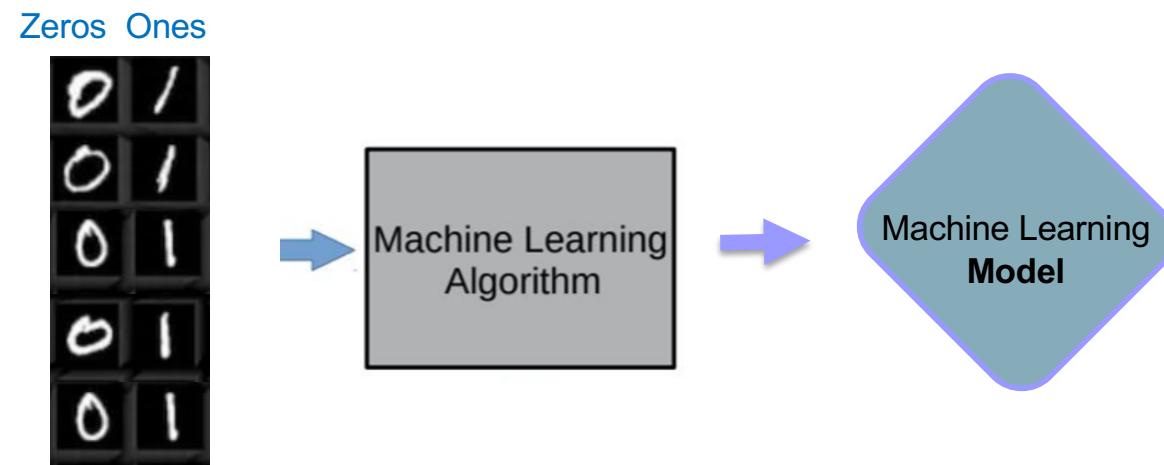
Black pixels are 0,
White pixels are 255

$$\begin{matrix} \boxed{I} \\ \equiv \\ \begin{matrix} 0 \\ 0 \\ 0 \\ \dots \\ 1 \\ 1 \\ 1 \\ \dots \\ 1 \\ 0 \\ 0 \\ \dots \end{matrix} \end{matrix} \quad \text{Flattened Image (in column order)}$$

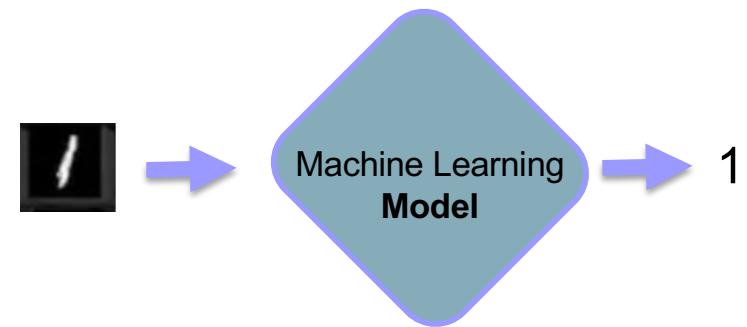
Simplified MNIST Dataset

Task: Classify a given digit image as 0 or 1

Learning: Using **labelled examples**, use a machine learning algorithm, to develop an ML model



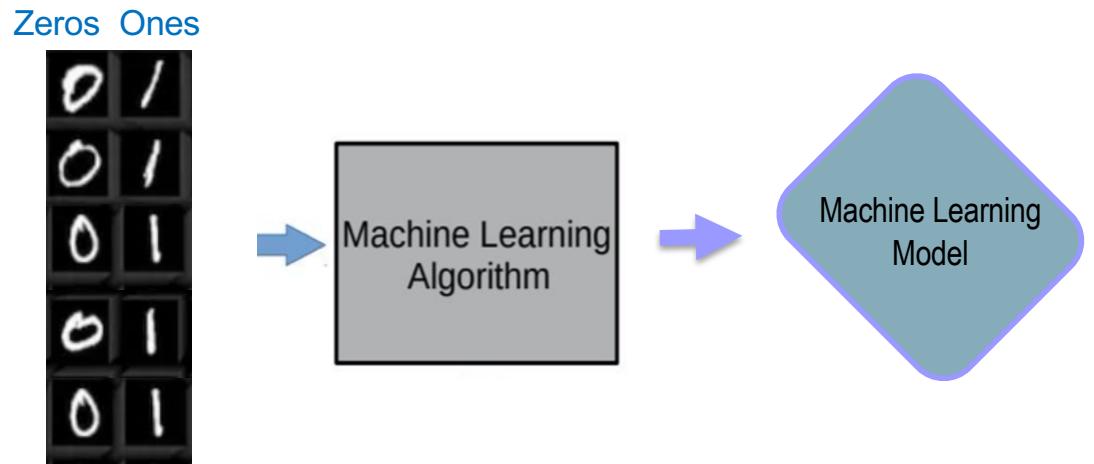
Query/Inference/Test: Given an **unlabelled example**, use the **model** to **classify** it



Simplified MNIST Dataset

Task: Classify the given digit image as 0 or 1

Learning: Considering labelled examples, use a ML algorithm, to develop an ML model



ML Model is a simple Decision Tree:

//If Center is Black

If sum(Pixel(54),Pixel(55), Pixel(56)) < 50

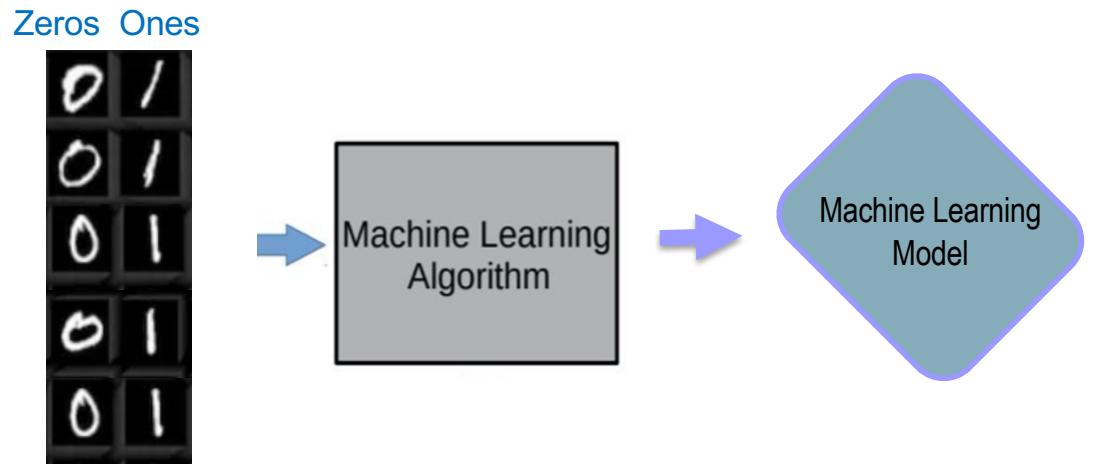
Then y=0; //Predict zero

Else y = 1; //Predict 1

Simplified MNIST Dataset

Task: Classify the given digit image as 0 or 1

Learning: Considering labelled examples, use a ML algorithm, to develop an ML model



ML Model is a simple Decision Tree:

//If Center is Black

If sum(Pixel(54),Pixel(55), Pixel(56)) < 50

 Then y=0; //Predict zero

 Else y = 1; //Predict 1

How do we know which pixels to look at?

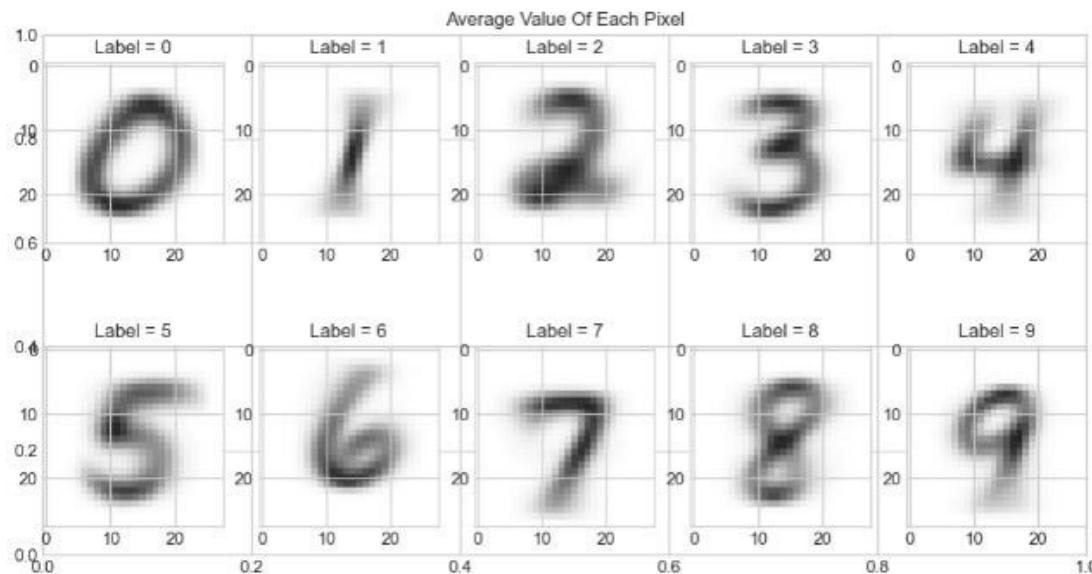
- Manually features (mostly in the past)
- Automatically learned features

Feature Extraction – Input Representation

- We could also **extract features** from the **raw image**

Manual feature extraction

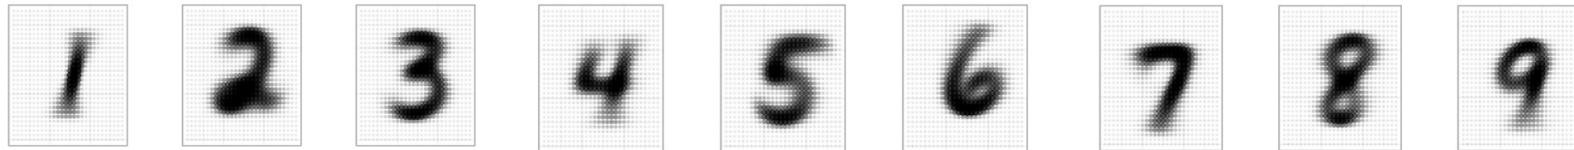
E.g. pixel averages to gain some invariance against shifts



When we talk about **features**, it will mean whatever features are deemed useful, and thus input into the system – raw or extracted, unless specified.

Learning w. Prototypes – A ML Model for MNIST

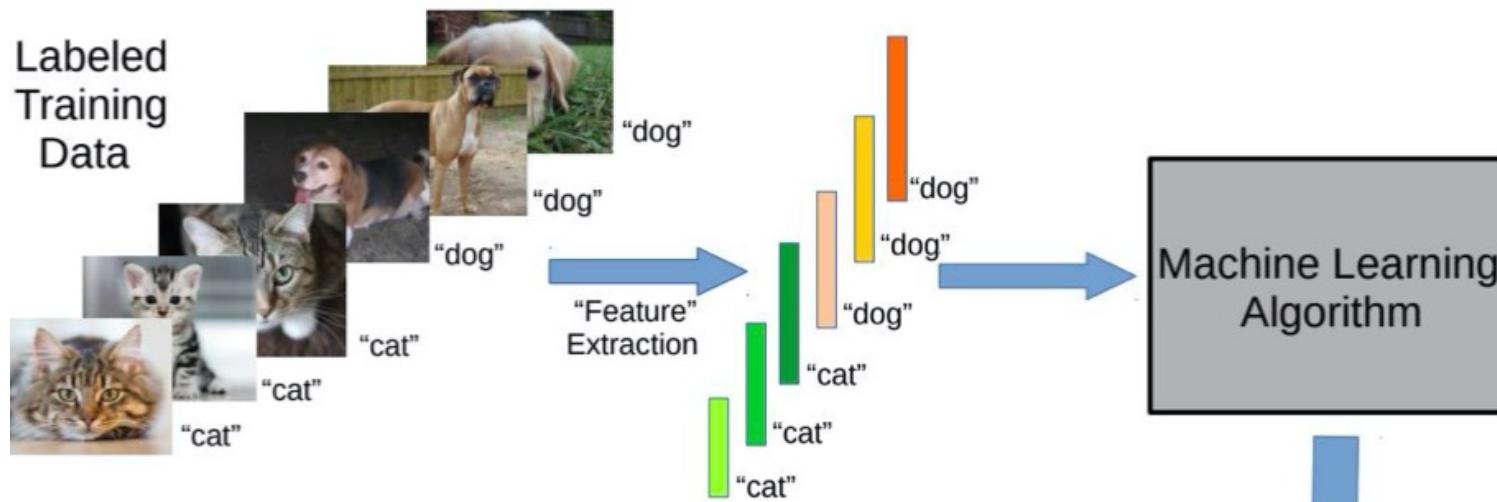
- **Basic idea:** Represent each class by a “prototype” vector
- **Class Prototype:** The “mean” or “average” of inputs from that



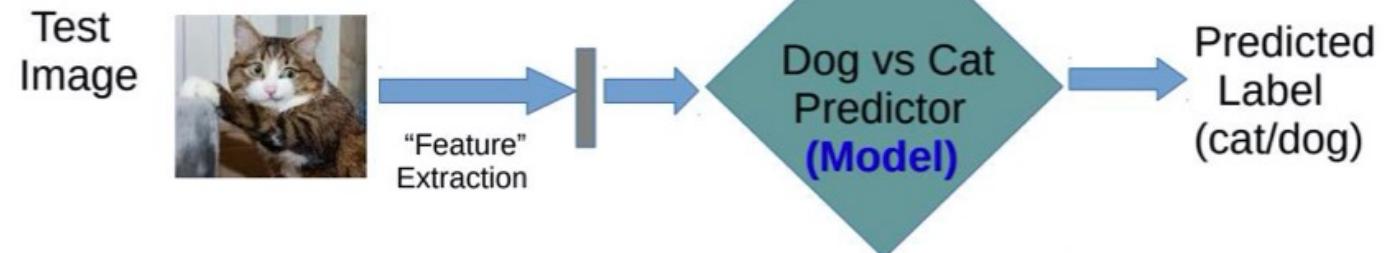
Averages (prototypes) of each of the handwritten digits 1-9

- **Inference:** Predict label of a test input based on its distances from the class prototypes
 - Predicted label will be the class that is the closest to the test input
- **Distance Metric:** Euclidean, weight Euclidean, Mahalanobis, ...

Cat vs Dog Binary Classification

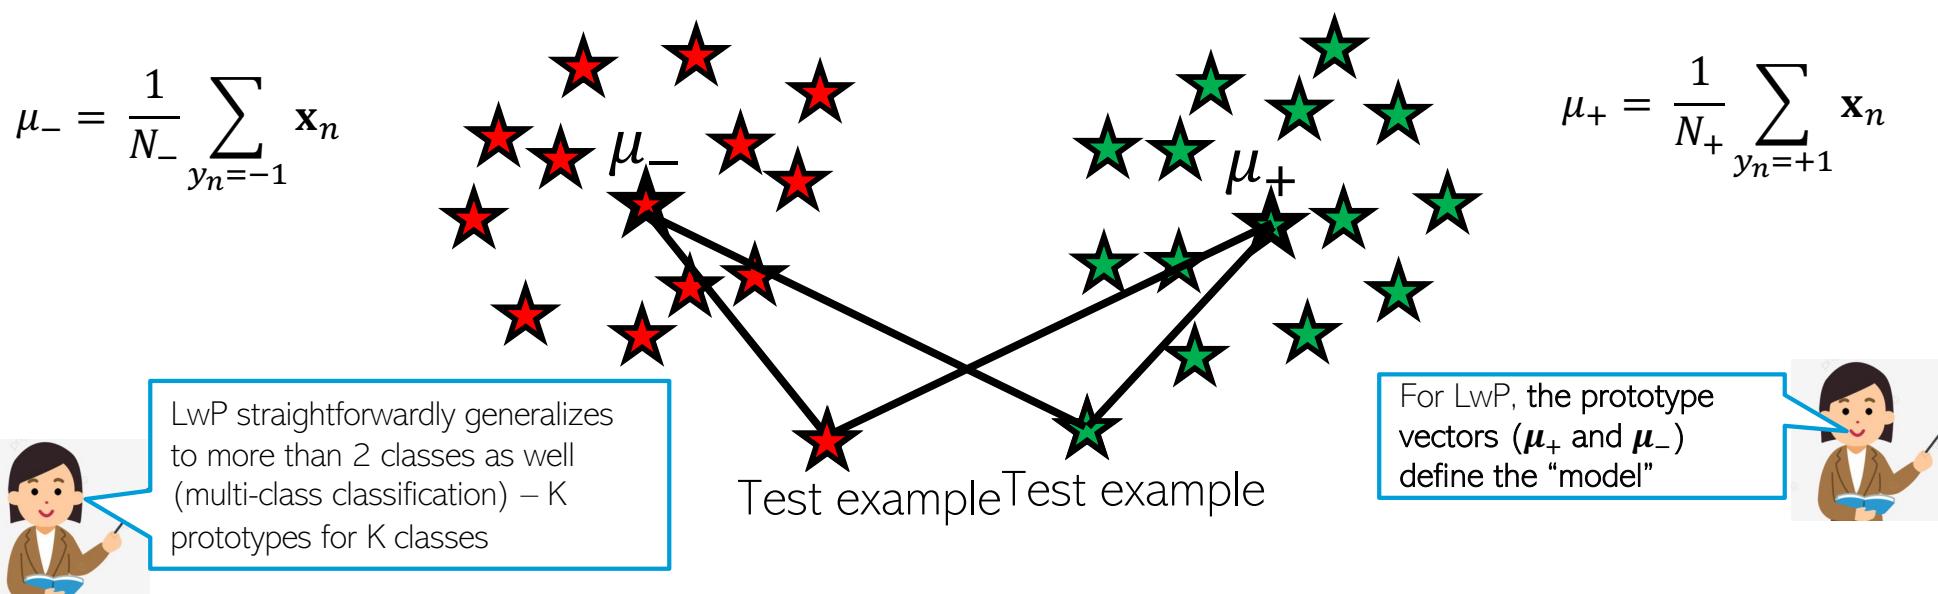


Note: The feature extraction phase may be part of the machine learning algorithm itself
(referred to "feature learning" or "representation learning")



Learning with Prototypes / Template Matching

- Suppose the task is binary classification (two classes assumed pos and neg)
- Training data: N labelled examples $\{(\mathbf{x}_n, y_n)\}_{n=1}^N, \mathbf{x}_n \in \mathbb{R}^D, y_n \in \{-1, +1\}$
 - Assume N_+ example from positive class, N_- examples from negative class
 - Assume green is positive and red is negative

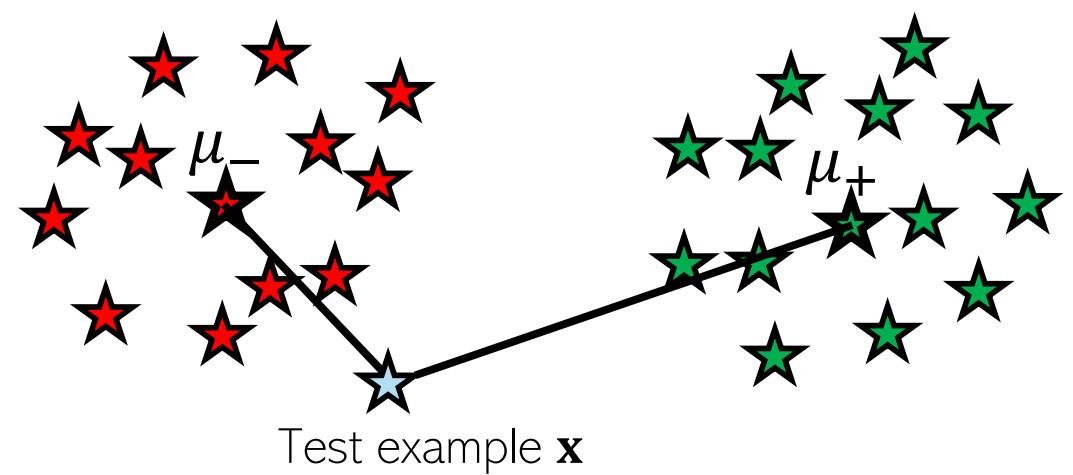


LwP: The Prediction Rule, Mathematically

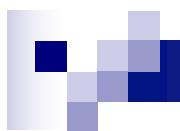
- What does the prediction rule for LwP look like mathematically?
- Assume we are using Euclidean distances here

$$\|\boldsymbol{\mu}_- - \mathbf{x}\|^2 = \|\boldsymbol{\mu}_-\|^2 + \|\mathbf{x}\|^2 - 2\langle \boldsymbol{\mu}_-, \mathbf{x} \rangle$$

$$\|\boldsymbol{\mu}_+ - \mathbf{x}\|^2 = \|\boldsymbol{\mu}_+\|^2 + \|\mathbf{x}\|^2 - 2\langle \boldsymbol{\mu}_+, \mathbf{x} \rangle$$

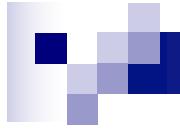


Prediction Rule: Predict label as +1 if $\|\boldsymbol{\mu}_- - \mathbf{x}\|^2 > \|\boldsymbol{\mu}_+ - \mathbf{x}\|^2$
otherwise as -1

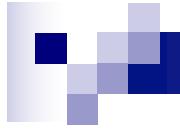


Template Matching Model - Key Aspects

- Very simple, interpretable, and lightweight model
 - Just requires computing and storing the class prototype vectors
- Works with any number of classes (thus for multi-class classification as well)
- Can be generalized in various ways to improve it further, e.g.,
 - Modeling each class by a probability distribution rather than just a prototype vector
 - Using distances other than the standard Euclidean distance (e.g., Mahalanobis)
- How well template matching works depends crucially on the distance metric. With a learned distance function, can work very well even with very few examples from each class (used in some “few-shot learning” models ...)

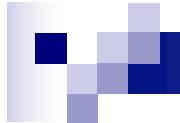


Performance Evaluation



How do we evaluate a trained model?

- We are not really interested in **how well a system does on the training data** but how well it will do in **general** (for **unseen** data)
- **Training set performance is mostly irrelevant**, but:
 - If too low, it tells us that our model's complexity may not be sufficient.
 - if too high, it doesn't tell us that we are doing well. The model may have even memorized all the labels



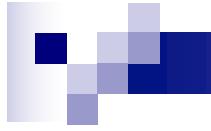
Empirical Risk Minimization

Assume we have learned a function $f(\mathbf{x})$, mapping the input \mathbf{x} to a label y

Formally, we want to **minimize the true risk** (also called **loss**):

$$R_{\text{true}}(f) = \mathbb{E}_p[L(f(\mathbf{x}), \mathbf{y})] = \int \int p(\mathbf{x}, \mathbf{y}) L(f(\mathbf{x}), \mathbf{y}) d\mathbf{x} d\mathbf{y}.$$

where $L(f(\mathbf{x}), y)$ is the loss incurred predicting $f(\mathbf{x})$ instead of the target value y ; and $p(\mathbf{x}, \mathbf{y})$ is the joint distribution of the input and labels.



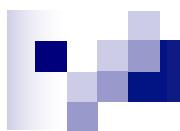
Empirical Risk Minimization

$$R_{\text{true}}(f) = \mathbb{E}_p [L(f(\mathbf{x}), \mathbf{y})] = \int \int p(\mathbf{x}, \mathbf{y}) L(f(\mathbf{x}), \mathbf{y}) d\mathbf{x} d\mathbf{y}.$$

- Since we don't know the distribution $p(\mathbf{x}, \mathbf{y})$, we will use **empirical risk minimization**.

We will minimize the error on an **unseen test set** hoping that it is a good enough **substitute for the true risk**.

$$R_{\text{emp}}(f) = \frac{1}{n} \sum_{i=1}^n \ell(f(x^{(i)}), y^{(i)})$$



Loss Function / Performance Measure

Performance of the learner can be measured in one of the following ways, as suitable for the application:

Accuracy:

- Misclassification rate (in classification problems)
- Mean Squared Error (in regression problems)
- General Loss functions (taking into account different costs for different mistakes)
- Precision, Recall, Area under the ROC Curve (will see later)
- ...

Training duration

Test duration

Model size

Performance Metrics

- Misclassification rate for Classification

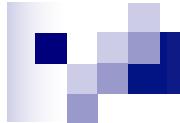
mistakes/ # test samples

$$\frac{1}{N} \sum_{i=1}^N I(t^i \neq y(x^i))$$

- Mean Squared Error for Regression

Average squared difference between the predicted value and the target values (we square so that errors don't cancel each other)

$$\frac{1}{N} \sum_{i=1}^N (y(x_i) - t_i)^2$$



Classification

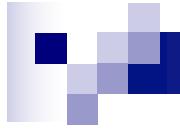
- Confusion matrix:

		<i>Predicted class</i>	
		Cancer	NotCancer
<i>Actual class</i>	Cancer	120	40
	NotCancer	50	450

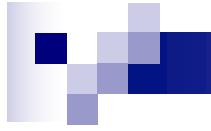
- Tells us where the errors are happening
- Applicable in multi-class problems as well

- Assume the rate of cancer is 5% among all people tested at a certain hospital.
- What happens if an ML model predicts all cases as non-cancer?

		<i>Predicted class</i>	
		Cancer	NotCancer
<i>Actual class</i>	Cancer	0	5
	NotCancer	0	95

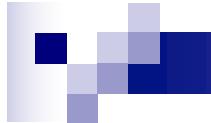


Validation and Cross-Validation



Train - Validation - Test

- If there is a need for parameter tuning (e.g. deciding how complex the decision rule/tree should be), we can use a **validation set**:
 - **Split the training data into two:**
 - validation (hold-out set)
 - development set
 - **Train different models with different parameters using samples from the development set**
 - **Measure performance them on the validation set to select the best parameters**
- The **test set** refers to the sequestered (kept away) data that is used as the indicator of generalization performance (e.g. Kaggle).



Cross-Validation

- If training set is not very large;
 - If the development (training portion) is kept larger, the learning improves
 - If the validation part is kept larger, generalization error estimate is more reliable
- For this reason, we often use cross-validation for model selection.
- You need to separate a subset as test, and another for validation.
 - Choose your model by testing on validation set
 - Test on test set.

k-fold cross-validation

Split the training set into k folds (random or **stratified**)

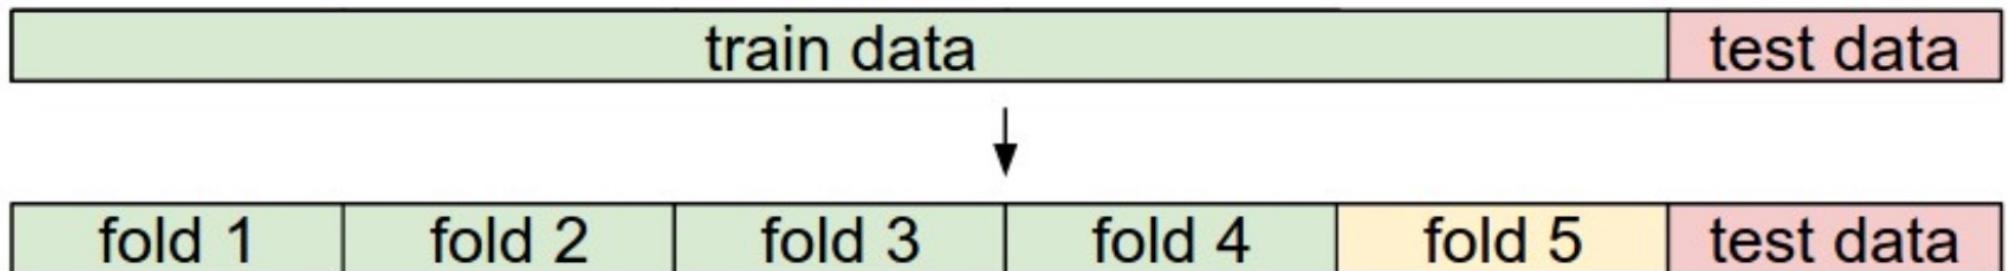
For each fold f

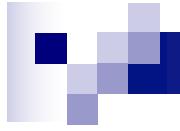
Train the learner with the **other** folds

Test the learner on the f fold, giving err_f

Average error = Average(err_f)

all f



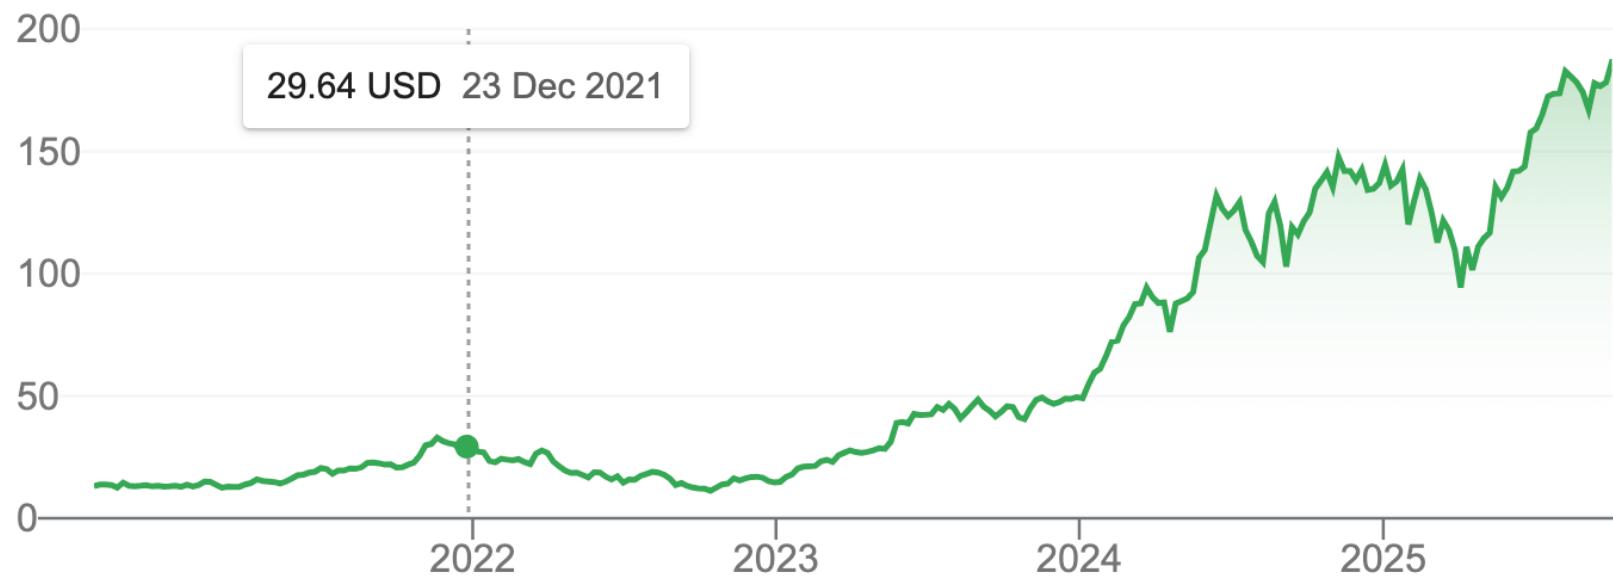


Test Set

- If you have access to the test data during training, you are not supposed to use it or even peek at it during your model building; **or your test performance becomes biased!**

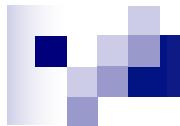
Stock Prediction

- Will NVIDIA's stock rise or fall tomorrow?



- It is impossible to reliably predict tomorrow`s stock price
 - Yet, investment companies do exactly that
 - We can detect patterns and predict the future probabilistically

- We need relevant factors as input
 - What if some relevant features are missing??
 - What if some irrelevant features are included??
- Yet, there could be a war tomorrow that changes everything
- Machine learning aims to detect historical patterns and behaviour so as to predict the future
 - Better than random? better than others? ...



!What to Know!

- Classification vs Regression
- Features vs Label
- Learning, inference...
- Machine learning algorithm vs ML model
- Raw input vs feature representation
- The MNIST and toy-MNIST dataset/problem
- Supervised learning terminology: $\{(\mathbf{x}_i, y_i)\}$
- Basic performance metrics, true vs empirical risk, misclassification rate, MSE, confusion matrix
- Prototype learning
 - >> Later: k-nearest neighbor model
- Correct use of validation, development and test sets