


Machine Learning

K-Nearest Neighbor (kNN)

Many slides/pictures thanks to Öznur Taştan Okan\ others as noted.



Nearest Neighbor Classifiers

- Basic idea:
 - Label a given test sample according to the labels of its k nearest neighbors.

1. Compute distance from the test sample to all observations
 2. Find k of the “nearest” samples
 3. Classify according to the majority label of these nearest neighbors


- ### - Special case with $k=1$

Training set

Test sample



00000000000000000000
11111111111111111111
22222222222222222222
33333333333333333333
44444444444444444444
55555555555555555555
66666666666666666666
77777777777777777777
88888888888888888888
99999999999999999999



Nearest ($k=1$) Neighbor Algorithm

- The training data $D = \{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)\}$
 - **Learning step:**
 - Store training examples ☺
 - Called a **lazy algorithm** for this reason
 - **Prediction step:**
 - Classify a new example \mathbf{x} by finding the training example (\mathbf{x}_i, y_i) that is nearest to \mathbf{x}
 - Predict the class y as y_i

K-Nearest Neighbour(kNN)





Photo by Sarang Anil Gotke on [Kdnuggets](#)

K-Nearest Neighbor (k-NN)

Formally:


- Given training data $D=\{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)\}$ and a test sample \mathbf{x}
- Find the **k** most-similar training examples
 - For classification, choose the **majority label** among those neighbors and assign to the test sample
 - For regression, choose the **average target** among those neighbors and assign to the test sample
- Algorithm requires meta-parameter **k** and a **distance function** to compute similarities between \mathbf{x} and training samples
- Special case: 1-NN Nearest neighbor



Important Decisions



- Value of k
- Distance measure
- Voting mechanism

Classification – Decision Boundary





k-NN Decision Boundary

- How does the decision boundary changes with k?
 - Which one is k=1?
 - Which one is k=7?




k-NN Decision Boundary

- How does the decision boundary changes with k?
 - Which ones is k=1?
 - Which one is k=7?



k-NN Decision Boundary

- Increasing k smooths the decision boundary
 - Smoother predictions, since we average over more data
 - Majority voting means less emphasis on individual points
- But could also be too smooth.
 - Think of the extreme case $k = N$, where N is number of training examples. What happens then?




Effect of k

- Very large value of k
 - Everything is classified as the most probable class
- Small value
 - Highly variable, small change to the training data will lead to large changes in the classifier


1- NN Decision Boundary

- The decision boundaries form a subset of the Voronoi diagram for the training data.
- Nearest-neighbor classifier produces **piecewise linear** decision boundaries




Choice of K


- Small K
 - Creates many small regions for each class
 - May lead to non-smooth decision boundaries and overfit
- Large K
 - Creates fewer larger regions
 - Usually leads to smoother decision boundaries (caution: too smooth decision boundary can underfit)
- Choosing K
 - Often data dependent and heuristic based
 - Or using cross-validation (using some held-out data)
 - In general, a K that is too small or too big is bad!



Selecting k


- Set aside a portion of your training data (validation set)
- Vary k, observe validation set error
- Pick k that gives best generalization performance as measured on the validation set






IMPORTANT DECISIONS

- Value of k (usually odd)
- Distance measure
- Voting mechanism




Distance Measures

- Key component of k-NN algorithm
 - Defines which examples are similar and which are not
 - Can have strong effect on performance



Distance - $d(x,y)$

- Numerical measure of how different (dissimilar) two data samples x and y are
 - A function that maps pairs of objects to real values
 - Lower when objects are more alike
- Minimum distance is 0, when comparing an object with itself.



Distance Metric

- A distance function d is a **distance metric** if it is a function from pairs of objects to real numbers such that:
 1. $d(x,y) \geq 0$ (**non-negativity**)
 2. $d(x,y) = 0$ iff $x = y$ (**identity**)
 3. $d(x,y) = d(y,x)$ (**symmetry**)
 4. $d(x,y) \leq d(x,z) + d(z,y)$ (**triangle inequality**)

Distances for Real Vectors

- Euclidean distance:

$$dist(x, y) = \sqrt{\sum_{i=1} (x_i - y_i)^2}$$

- Manhattan distance: L_p norm or Minkowski distance:

$$dist(x, y) = \sum_{i=1} |x_i - y_i|$$

$$dist(x, y) = \left(\sum_{i=1} |x_i - y_i|^p \right)^{1/p}$$

- The Euclidian distance between x and y is the L_2 -norm of $\Delta = \mathbf{x} - \mathbf{y}$ (the difference vector)
- The Manhattan distance between x and y is the L_1 -norm of $\Delta = \mathbf{x} - \mathbf{y}$ (the difference vector)

Viewed as Vector Norms

- $\|\mathbf{x}\|_2 = \sqrt{\sum_i |x_i|^2}$


- $\|\mathbf{x}\|_1 = \sum_i |x_i|$

- $\|\mathbf{x}\|_p = \sqrt[p]{\sum_i |x_i|^p}$

Example of Distances

L_2 -norm:

$$dist(x,y) = \sqrt{4^2 + 3^2} = \sqrt{25}$$




L_1 -norm:

$$dist(x,y) = 4 + 3 = 7$$

L_∞ -norm:

$$dist(x,y) = \max\{3, 4\} = 4$$



Distance Between Strings

- How do we define similarity between strings?

weird	wierd
intelligent	unintelligent
Athena	Athina

- Important for recognizing and correcting typing errors and analyzing DNA sequences

Hamming Distance


- Hamming distance is the number of positions in which bit-vectors differ.

- Example: $p_1 = 10\textcolor{red}{1}01$
 $p_2 = 10\textcolor{red}{0}11$

$d(p_1, p_2) = 2$ because the bit-vectors differ in the 3rd and 4th positions.

- Hamming distance between two vectors of categorical attributes is the number of positions in which they differ.

- Example: $p_1 = (\text{red}, \text{tall}, \text{heavy})$,
 $p_2 = (\text{green}, \text{tall}, \text{light})$
 $d(p_1, p_2) = 2$



Scale Issues

- Attributes may have to be scaled/normalized to prevent distance measures from being dominated by one of the attributes


- Example:
 - height of a person may vary from 1.5 m to 1.8 m
 - weight of a person may vary from 40kg to 200kg
 - income of a person may vary from 1,000 to 10,000TL

Scale Issues

- For real-valued feature vectors, we can use Euclidean distance

$$D(u, v)^2 = \|u - v\|^2 = (u - v)^T (u - v) = \sum_{i=1}^d (u_i - v_i)^2$$

- If we scale x_1 by 1/3, NN changes!



Scaling/Normalizing Features

For a feature $x \in \mathcal{R}$:

- **Linear scaling to unit range**

– Given a lower and upper bound, $[a,b]$


$$\tilde{x} = \frac{x - a}{b - a}$$

- **Linear scaling to unit variance**

– Transform to zero mean and unit variance as


$$\tilde{x} = \frac{x - \mu}{\sigma}$$

where μ is the sample mean and σ is the sample standard deviation of the feature.



IMPORTANT DECISIONS

- Value of k (usually odd)
- Distance measure
- Voting mechanism



Combining of Neighbor Labels


- Options for determining the class from nearest neighbor list
 - Take **majority** vote of class labels among the k -nearest neighbors
 - **Weight** the votes according to distance
 - example: weight factor $w = 1 / d^2$

Summary

- Pros:
 - **Simple and intuitive; easily implementable**
 - *Asymptotically consistent:*
 - **With infinite training data and large enough K, K-NN approaches the best possible classifier ("Bayes optimal")**
- Cons:
 - **Store all the training data in memory even at test time**
 - Can be memory intensive for large training datasets
 - An example of **non-parametric**, or **memory/instance-based methods** in contrast to parametric, model-based learning models
 - **Expensive at test time: $O(ND)$ computations for each test point!**
 - Have to search through all training data to find nearest neighbors
 - Distance computations with N training points (D features each)
 - **Sensitive to noisy features**
 - **May perform badly in high dimensions (curse of dimensionality)**
 - In high dimensions, distance notions can be counter-intuitive!



K-NN REGRESSION



k-NN for regression

- As in classification, we find the label of the k nearest data points
- But instead of using the majority label, we will take the average of the labels of the k neighbors

k-NN for regression

- Query is $x=4$.
 - When $k=1$, we find the point to the left as closest (around $x=3.2$). Thus, the predicted $y(4)=6$.
 - When $k=2$, we find the points to the left and right closest (around $x=3.2$ and $x=5$). Thus, the predicted $y(4)=\text{average}(6,3)$.
 - Thus, the 3 red dots are the predicted y values for $x=4$, when $k=1..3$

