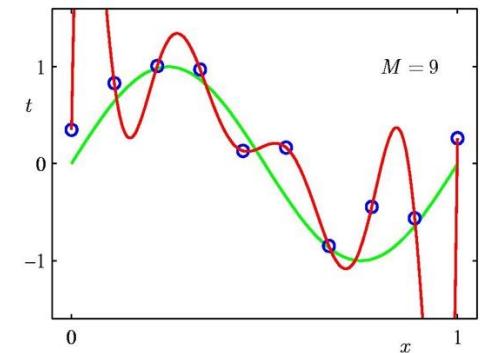
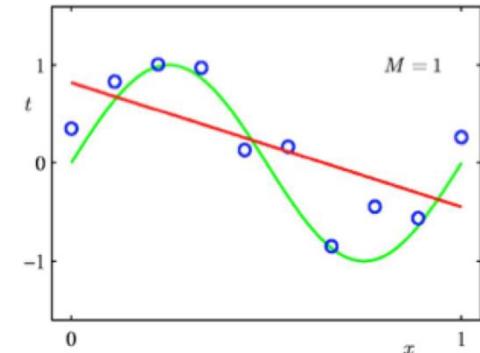


# Machine Learning Overfitting, Model Complexity, Regularization

Berrin Yanikoglu  
Sabancı Üniversitesi

# Overfitting and Model Complexity

- Imagine that we have some training data and we want to learn the underlying function between the independent variable  $x$  and the target values  $t$ .
- We can fit polynomials in varying degrees: lines to higher degree polynomials.
  - Higher degrees make the polynomial very capable to bend/flex to match the data as it has many parameters to change/adapt.
- However having zero train error does not mean the model (high order polyn.) will also have high generalization performance.
- In fact, a simpler model that has a similar performance compared to a more complex model, is often preferred (more on these later and more formally).

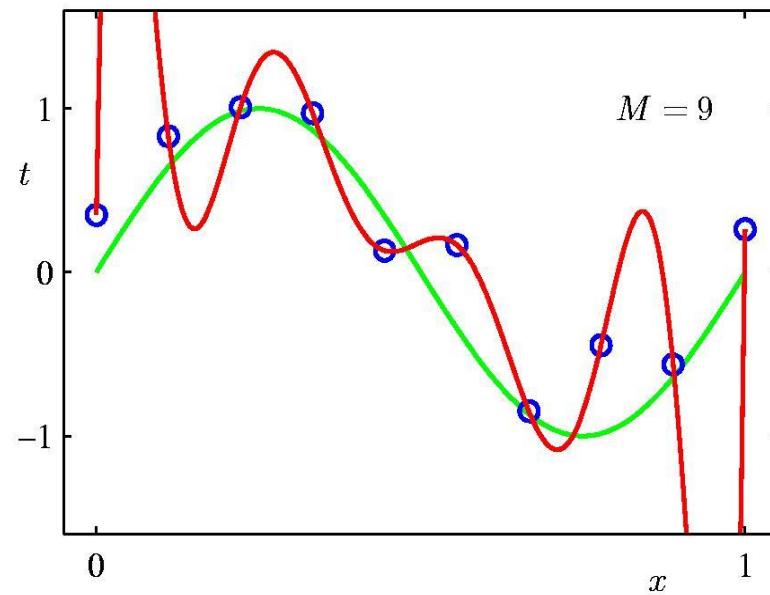
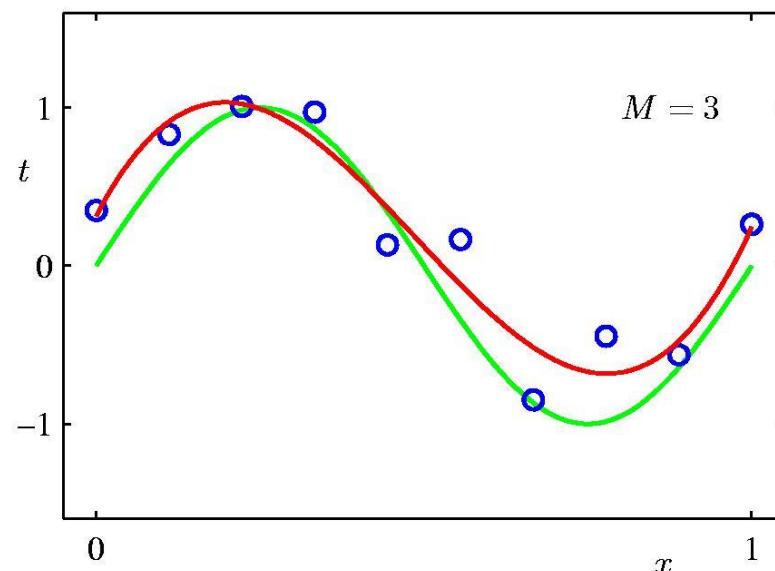
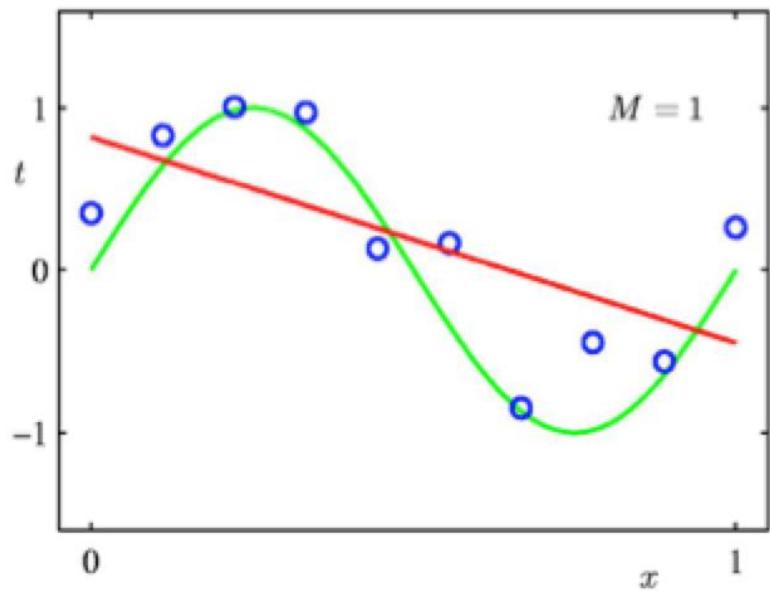


# 1<sup>st</sup> Order Polynomial

- Let's represent this function with polynomials of varying degree:

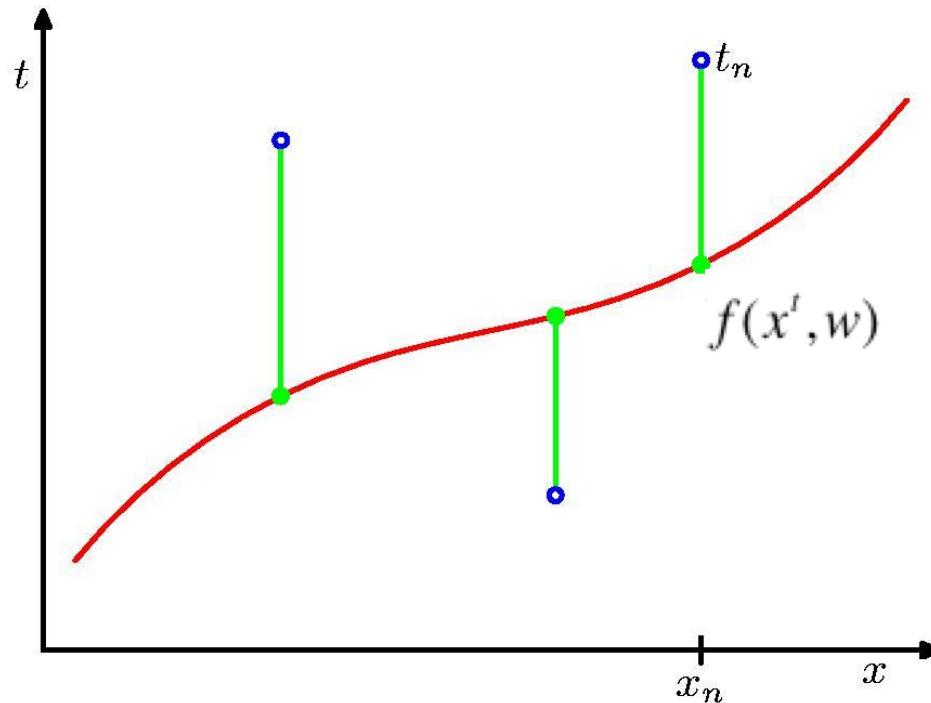
$$y(x) = w_0 + w_1x + w_2x^2 + \dots + w_Mx^M = \sum_{j=0}^M w_j x^j$$

# 1<sup>st</sup> Order Polynomial



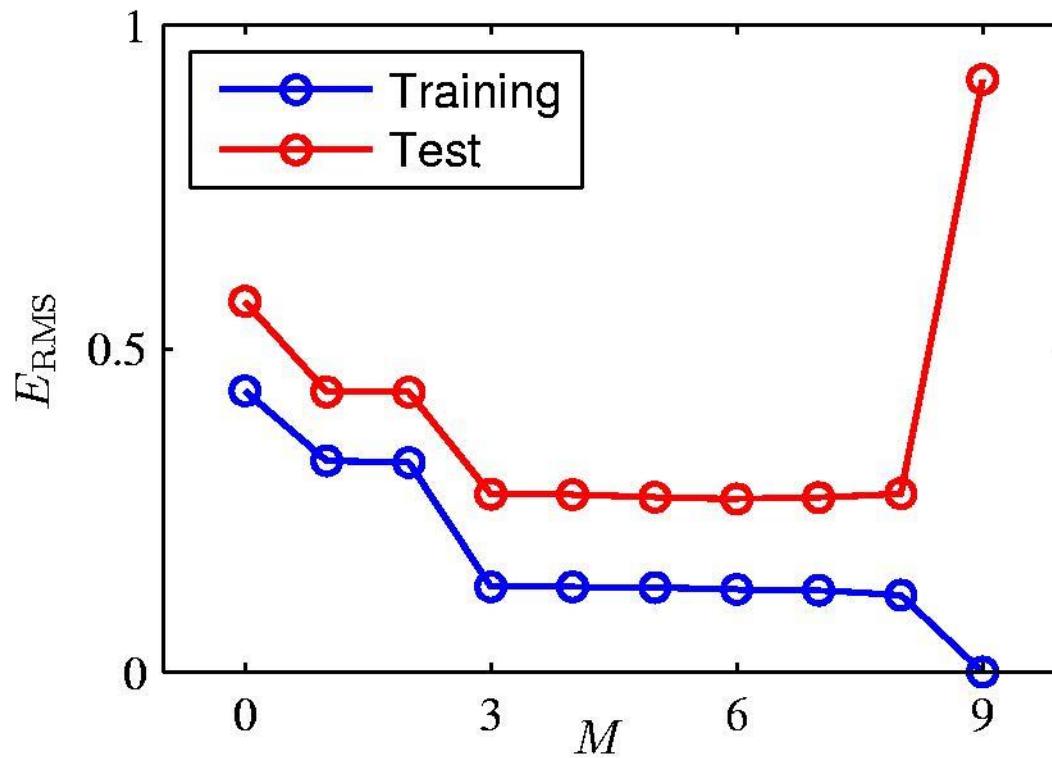
- We do not know yet which is the best model, maybe the 9<sup>th</sup> degree polynomial after all.
- The main/typical approach is to use the validation set performance to decide which model to choose.

# Sum-of-Squares Error Function



$$E(\mathbf{w}) = \frac{1}{2} \sum_{i=1}^N (y(x^i) - t^i)^2$$

# Over-fitting



Root-Mean-Square (RMS) Error:

$$E_{\text{RMS}} = \sqrt{2E(\mathbf{w}^*)/N}$$

# Polynomial Coefficients

	$M = 0$	$M = 1$	$M = 3$	$M = 9$
$w_0^*$	0.19	0.82	0.31	0.35
$w_1^*$		-1.27	7.99	232.37
$w_2^*$			-25.43	-5321.83
$w_3^*$			17.37	48568.31
$w_4^*$				-231639.30
$w_5^*$				640042.26
$w_6^*$				-1061800.52
$w_7^*$				1042400.18
$w_8^*$				-557682.99
$w_9^*$				125201.43

## Formal definition:

A hypothesis  $f$  is said to overfit the training data if there exists another hypothesis,  $f'$ , such that  $f$  has smaller error than  $f'$  on the training data, but  $f'$  has smaller error on the test data than  $f$ .

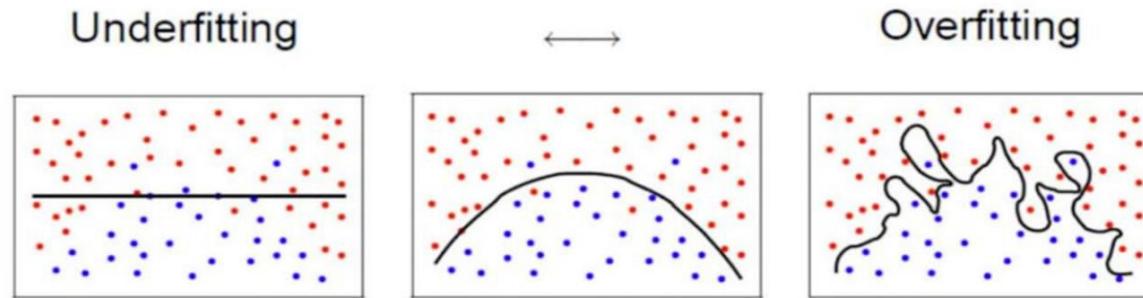
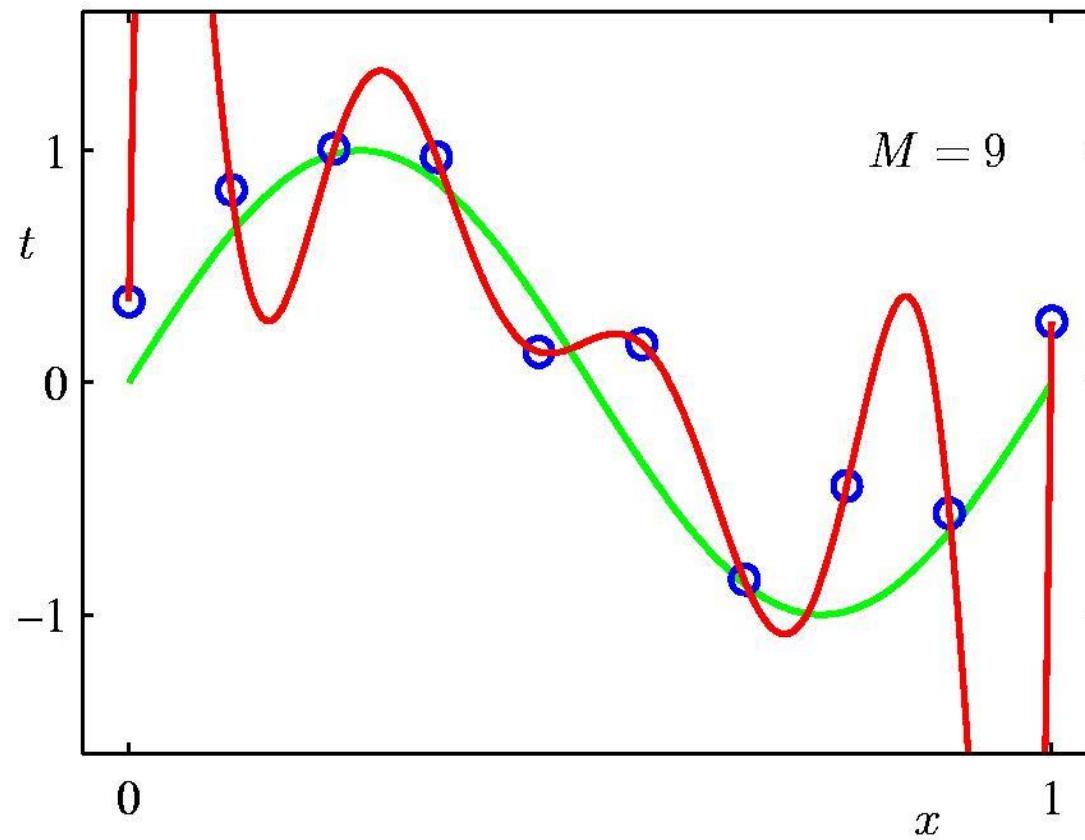


Image Credit: <https://tomrobertshaw.net/2015/12/introduction-to-machine-learning-with-naive-bayes>

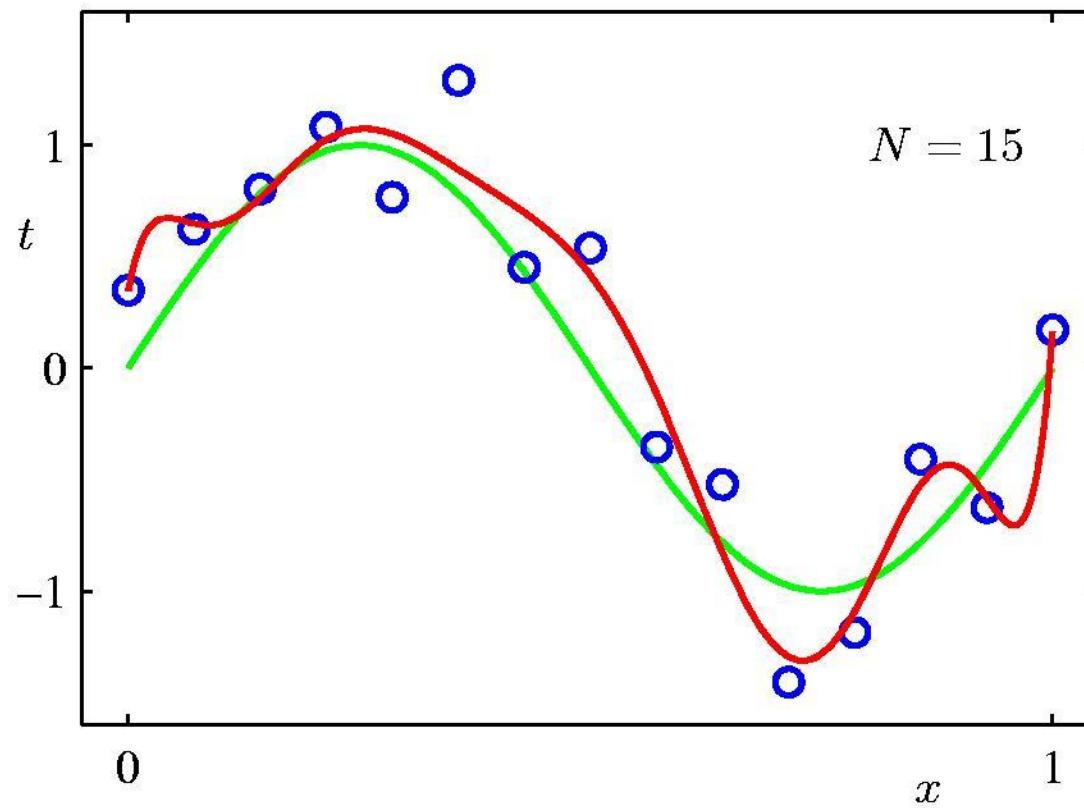


# Effect of Data

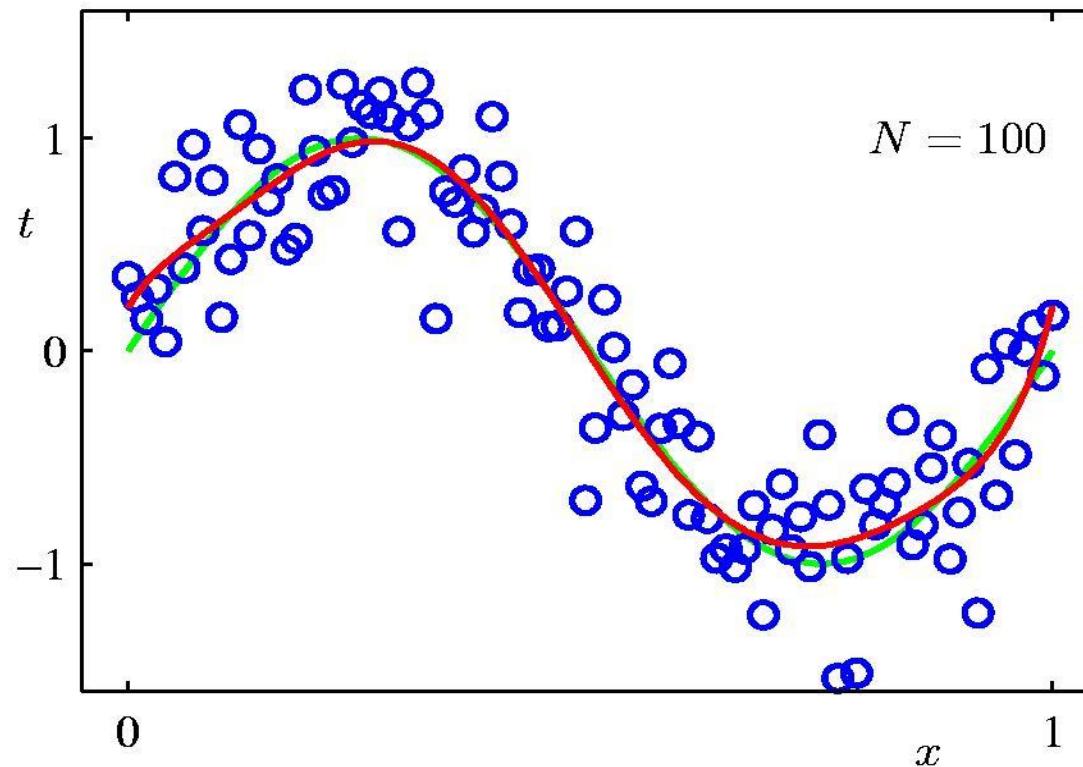
9<sup>th</sup> Order Polynomial



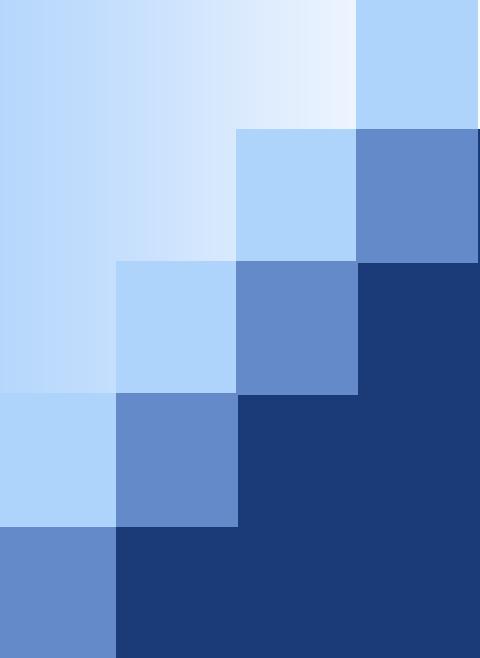
9<sup>th</sup> Order Polynomial



## 9<sup>th</sup> Order Polynomial



As data increases, the polynomial is further bound, reducing wild changes.



# Regularization

# Regularization

- Use complex models, but penalize large coefficient values:

$$\text{Min}_{\mathbf{w}, \mathbf{b}} (\text{MSE} + \text{penalty}) = \text{Min} \left[ \frac{1}{N} \sum_{i=1}^N (y_i - f_{\mathbf{w}, \mathbf{b}}(X_i))^2 + \text{penalty}(\mathbf{w}) \right]$$

Fit data                      Regularize

Lasso (L1 Regularization):

$$\text{minimize} \left( \frac{1}{N} \sum_{i=1}^N (y_i - f(\mathbf{x}_i; \mathbf{w}))^2 + \lambda \sum_{j=1}^p |w_j| \right)$$

Ridge (L2 Regularization):

$$\text{minimize} \left( \frac{1}{N} \sum_{i=1}^N (y_i - f(\mathbf{x}_i; \mathbf{w}))^2 + \lambda \sum_{j=1}^p w_j^2 \right)$$

# Regularization

- Use complex models, but penalize large coefficient values:

$$\text{Min}_{\mathbf{w}, \mathbf{b}} (\text{MSE} + \text{penalty}) = \text{Min} \left[ \frac{1}{N} \sum_{i=1}^N (y_i - f_{\mathbf{w}, \mathbf{b}}(X_i))^2 + \text{penalty}(\mathbf{w}) \right]$$

Fit data                      Regularize

Lasso (L1 Regularization):

$$\text{minimize} \left( \frac{1}{N} \sum_{i=1}^N (y_i - f(\mathbf{x}_i; \mathbf{w}))^2 + \lambda \sum_{j=1}^p |w_j| \right)$$

Ridge Regression (L2 Regularization)

$$\text{minimize} \left( \frac{1}{N} \sum_{i=1}^N (y_i - f(\mathbf{x}_i; \mathbf{w}))^2 + \lambda \sum_{j=1}^p w_j^2 \right)$$

The tuning parameter  $\lambda$  serves to control the relative impact of the penalty term on the regression coefficient estimates.

# Lasso Regression

## Lasso (L1 Regularization):

$$\text{minimize} \left( \frac{1}{N} \sum_{i=1}^N (y_i - f(\mathbf{x}_i; \mathbf{w}))^2 + \lambda \sum_{j=1}^p |w_j| \right)$$

- LASSO stands for “Least Absolute Shrinkage and Selection Operator”
- LASSO eliminates the least important features from the model, it automatically performs a type of **feature selection**.
- It is best to apply regularization after **variable standardization** (weights should be of comparable scale, as the L1 or L2 norm is minimized as a whole).
- $\lambda$  should be chosen using cross-validation

# Ridge Regression

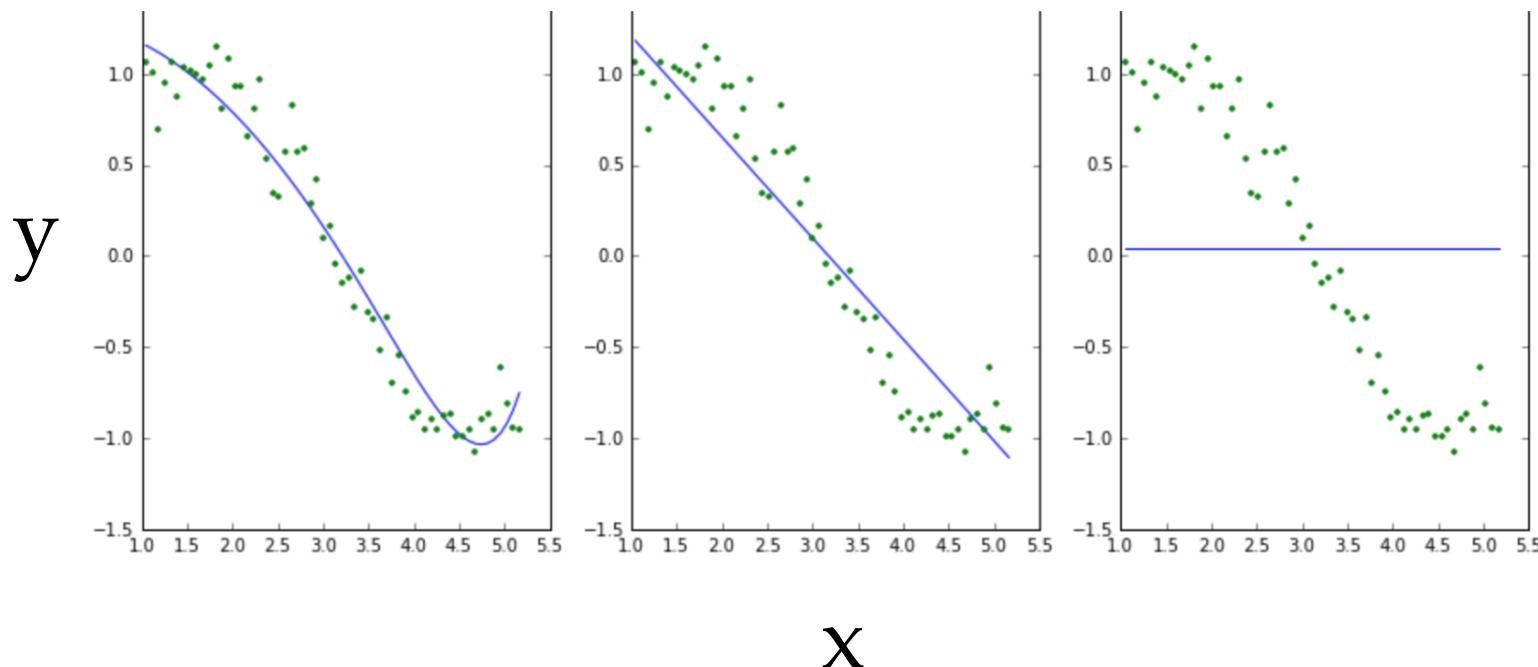
## Ridge Regression (L2 regularization)

$$\text{minimize} \left( \frac{1}{N} \sum_{i=1}^N (y_i - f(\mathbf{x}_i; \mathbf{w}))^2 + \lambda \sum_{j=1}^p w_j^2 \right)$$

- The shrinkage penalty has the effect of shrinking the estimates of  $w$  towards zero (not selecting one in case of multiple colinear features)
- It is best to apply regularization after **variable standardization** (weights should be of comparable scale, as the L1 or L2 norm is minimized as a whole).
- $\lambda$  should be chosen using cross-validation

# Effect of Lambda

- Blue line is the predicted function.
- Green dots are the training data
- As lambda increases the model becomes simpler.



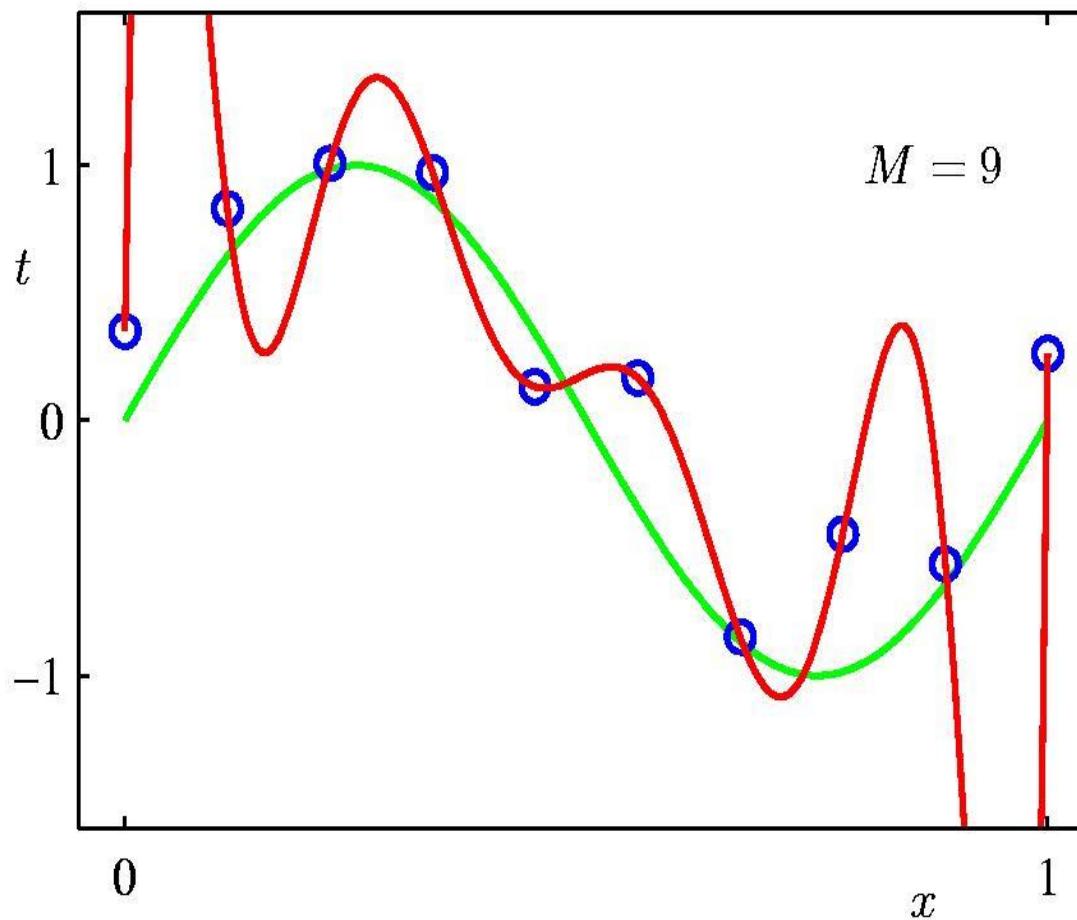
- When there are multiple colinear features, **L1 regularization (Lasso)** will often set the weight of one to zero. So, it is often used for **feature selection and model interpretability** (you may ignore features with zero weights).
- **L2 regularization (Ridge)** will decrease the weights of two colinear features, w/o setting one to zero, so not used for feature selection or simplicity, but may be better in terms of generalization.

# Regularization on 9<sup>th</sup> Order Polynomial

Example from the  
Bishop book.

$$\ln \lambda = -\infty$$

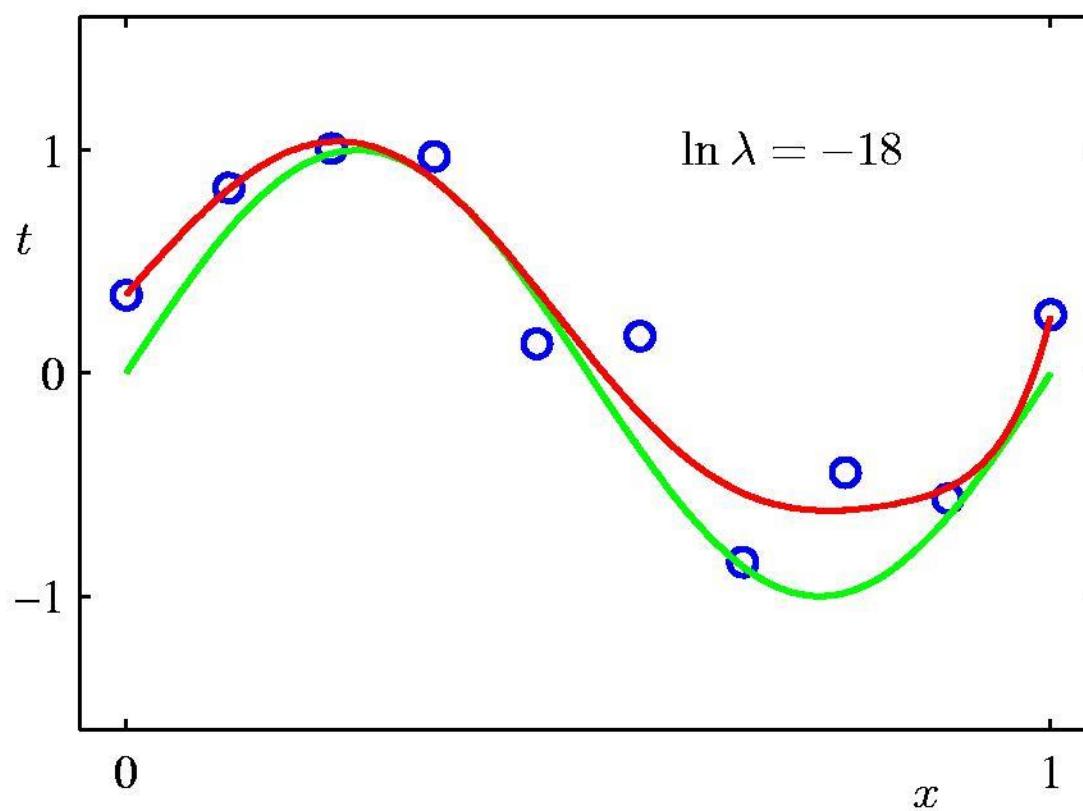
Too small  $\lambda$  – no regularization effect



# Regularization on 9<sup>th</sup> degree polynomial:

$$\ln \lambda = -18$$

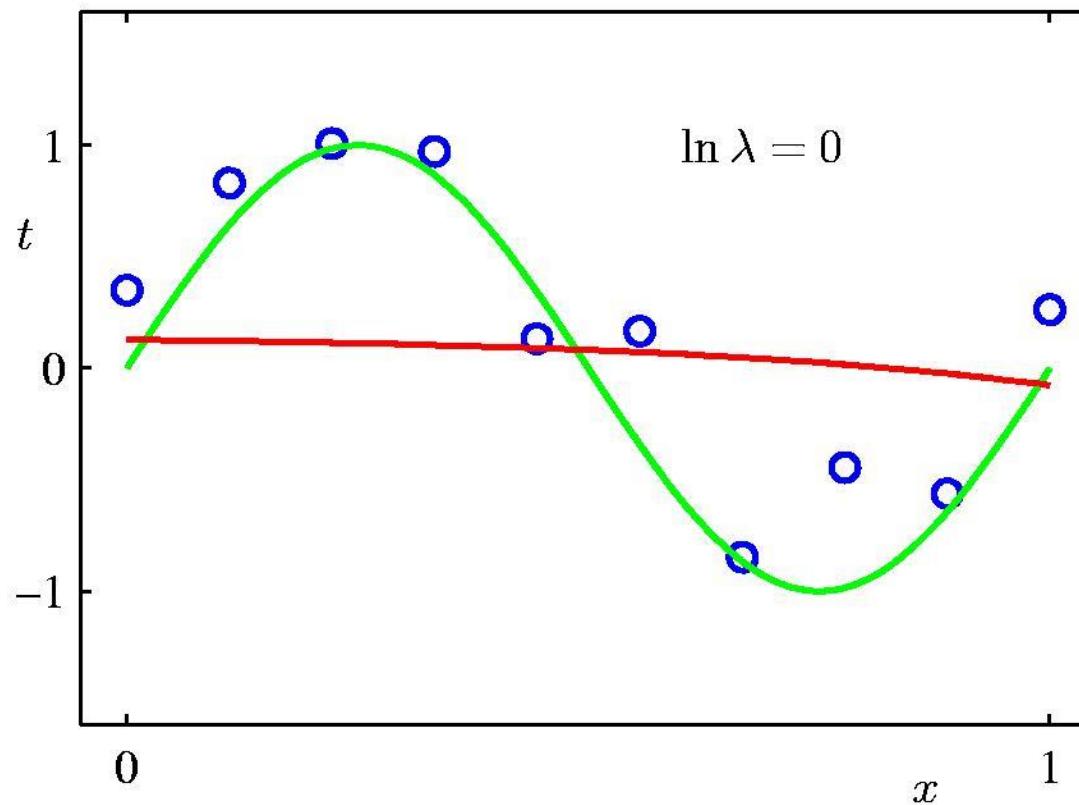
Right  $\lambda$  –good fit



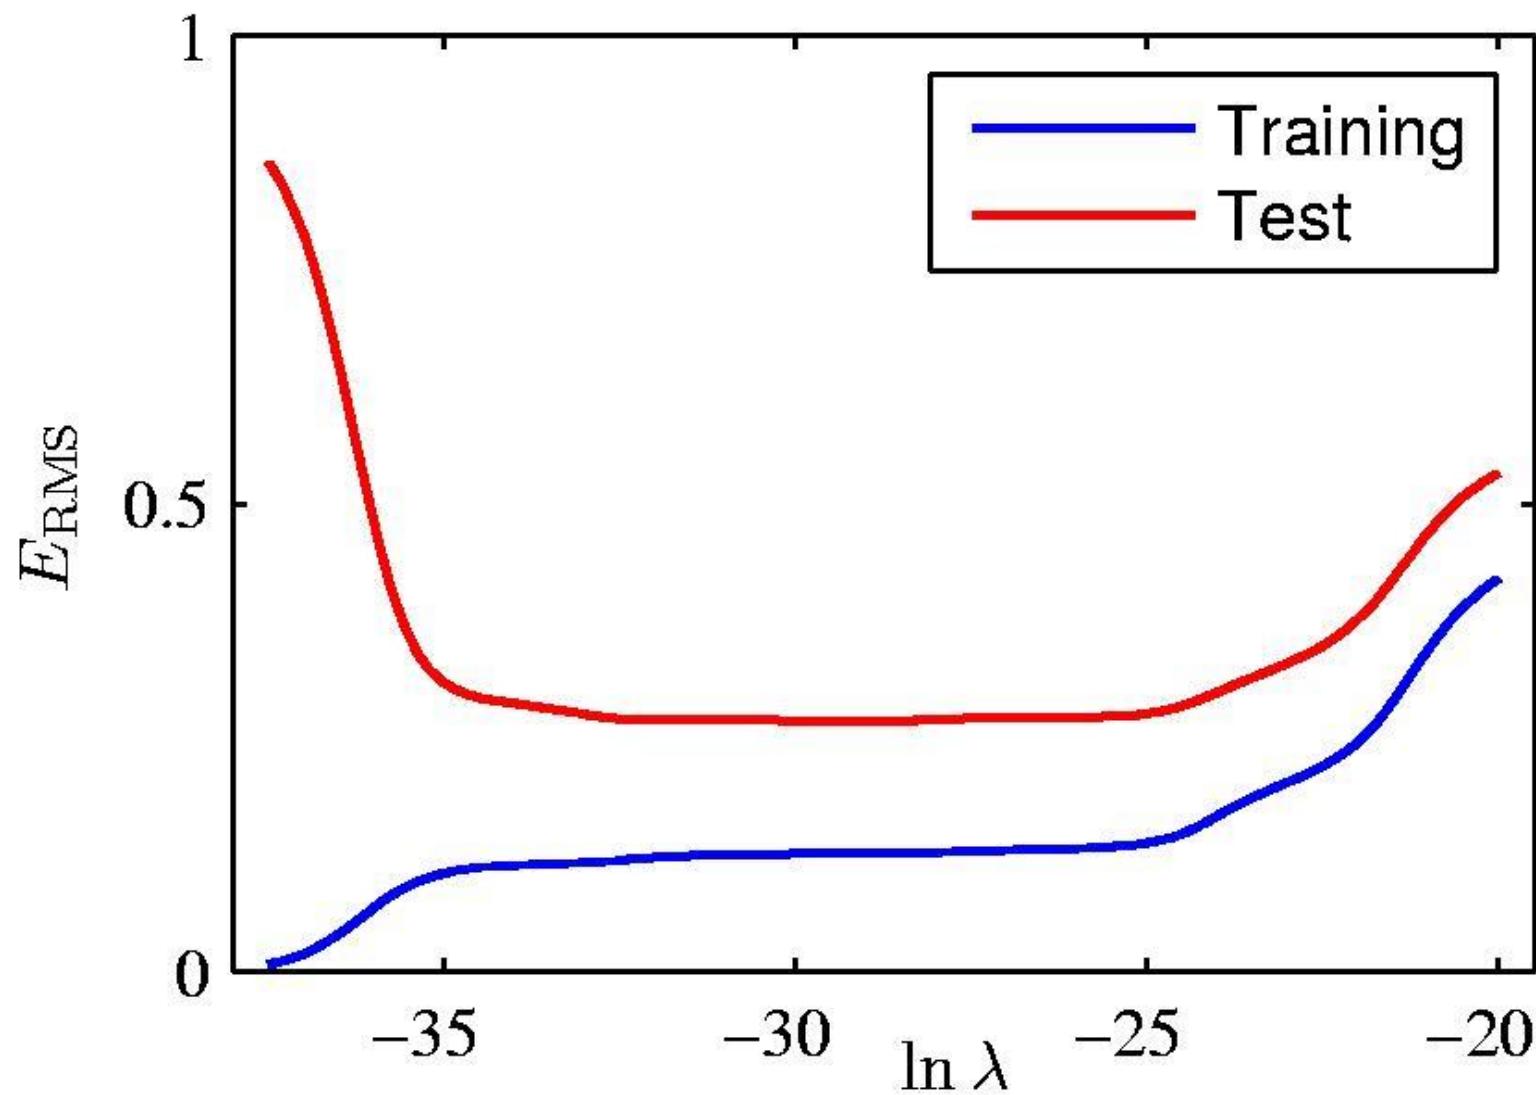
# Regularization:

$$\ln \lambda = 0$$

Large  $\lambda$  –regularization dominates



## Regularization: $E_{\text{RMS}}$ vs. $\ln \lambda$



# Polynomial Coefficients

	$\ln \lambda = -\infty$	$\ln \lambda = -18$	$\ln \lambda = 0$
$w_0^*$	0.35	0.35	0.13
$w_1^*$	232.37	4.74	-0.05
$w_2^*$	-5321.83	-0.77	-0.06
$w_3^*$	48568.31	-31.97	-0.05
$w_4^*$	-231639.30	-3.89	-0.03
$w_5^*$	640042.26	55.28	-0.02
$w_6^*$	-1061800.52	41.32	-0.01
$w_7^*$	1042400.18	-45.95	-0.00
$w_8^*$	-557682.99	-91.53	0.00
$w_9^*$	125201.43	72.68	0.01

Very little

Too much

Regularization

How to select the polynomial degree and regularization coefficient?

## Grid search:

for degree d in [1, 2, 3]:

//inner loop – compute val error for each  $\lambda$

for  $\lambda$  in [0, 0.1, 1, 10, ...]:

train model(degree,  $\lambda$ ) on training set

$\text{val\_error}(d, \lambda)$  = validation error of the model

//select best  $\lambda$  for degree

$\text{min\_val\_error}(d) = \min_{\lambda} \text{val\_error}(d, \lambda)$

//Select best model

$\text{min\_val\_error} = \min_d \text{val\_error}(d)$



# What You Should Know

- Different regression types (linear, polynomial, multiple linear)
- Least Squares and Gradient Descent solutions to regression
- When you may choose one or the other or how to choose the right complexity
  - You may choose a linear model if you have some prior expectation for the linearity
  - If you have large amounts of data, high complexity models may not pose a problem
- Regularization
  - You should use regularization to control your model complexity
- Final model selection should be done via grid search over hyperparameters (degree and lambda)