

Stochastic/Approximate Gradient Descent

Approximate Gradient Descent (Stochastic Backpropagation)

Normally, in gradient descent, we would need to compute how the error over all input samples (true gradient) changes with respect to a small change in a given weight.

But the common form of the gradient descent algorithm takes one input pattern, **compute the error of the network on that pattern only**, and updates the weights using only that information.

- Notice that the new weight may not be good/better for all patterns, but we expect that if we take a small step, we will average and approximate the true gradient.

Stochastic Approximation to Steepest Descent

Instead of updating every weight until all examples have been observed, **we update on every example:**

$$\nabla w_i \approx \eta (t-o) x_i$$

Remarks:

- Speeds up learning significantly when data sets are large
 - **Use a smaller learning step!**
- When there are multiple local minima, stochastic gradient descent may avoid the problem of getting stuck on a local minimum.

Skipped here

Gradient Descent Backpropagation Algorithm

Derivation for
General Activation Functions Networks

Stochastic Backpropagation

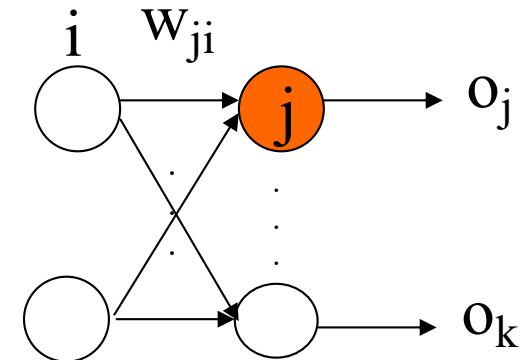
To calculate the partial derivative of E (here: loss on a single input) w.r.t a given weight w_{ji} of a node j , we have to consider whether this is the weight of an output or hidden node:

If w_{ji} is an **output** node weight, the situation is simple.

We always use the **Chain Rule**

Assuming multi-label MSE loss (simplest to show):

$$E = \sum_{j=1}^K (t_j - o_j)^2$$



$$o_j = f(\text{net}_j)$$

$$\text{net}_j = \sum_i o_i w_{ji}$$

Stochastic Backpropagation

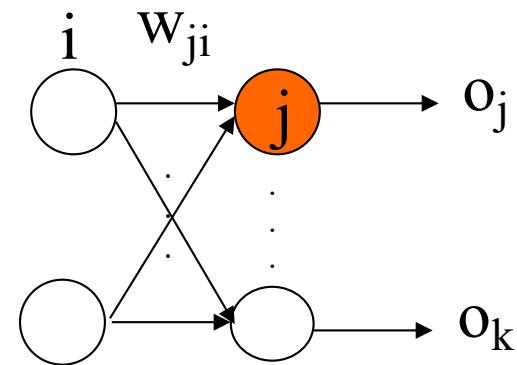
To calculate the partial derivative of E w.r.t a given weight w_{ji} of a node j , we have to consider whether this is the weight of an output or hidden node:

If w_{ji} is an **output** node weight, the situation is simple. We use the **Chain Rule**:

$$\frac{dE}{dw_{ji}} = \frac{dE}{do_j} \times \frac{do_j}{dnet_j} \times \frac{dnet_j}{dw_{ji}}$$

$$\frac{dE}{dw_{ji}} = -(t_j - o_j) \times f'(net_j) \times o_i$$

Note that output of node i (o_i) is the input to node j .



$$E = \sum_{j=1}^K (t_j - o_j)^2$$

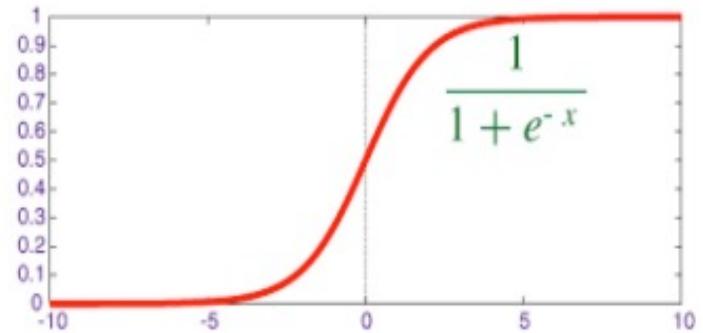
$$o_j = f(net_j)$$

$$net_j = \sum_i o_i w_{ji}$$

Transfer Function Derivatives

Sigmoid:

$$\begin{aligned} f'(n) &= \frac{d}{dn} \left(\frac{1}{1 + e^{-n}} \right) = \frac{e^{-n}}{(1 + e^{-n})^2} \\ &= \left(1 - \frac{1}{1 + e^{-n}} \right) \left(\frac{1}{1 + e^{-n}} \right) = \boxed{\left(1 - \frac{1}{1 + e^{-n}} \right) \left(\frac{1}{1 + e^{-n}} \right)} \end{aligned}$$



Linear: $f'(n) = \frac{d}{dn}(n) = 1$

- Computing the derivative is **very easy after the forward pass**
- Sigmoid nodes **saturate when output is large in magnitude.**
- Earlier layers learn much slower – **vanishing gradient**

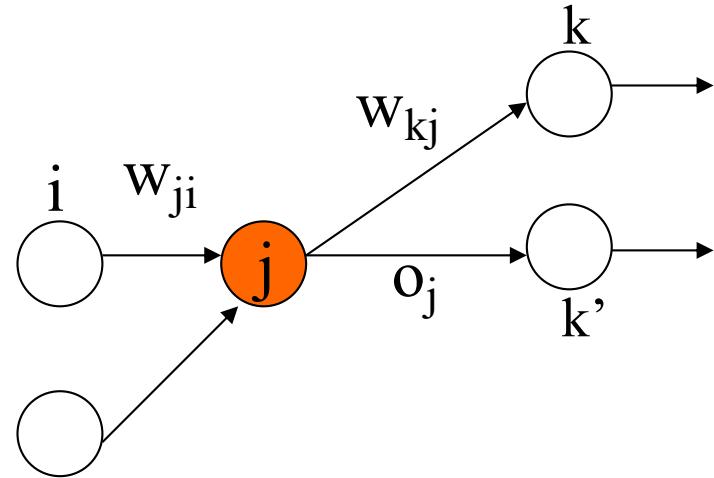
Backpropagation – Hidden nodes

The situation is more complex with a hidden node, because we don't know what the output of a hidden node should be how it affects loss.

If w_{ji} is a **hidden** node weight:

$$\frac{dE}{dw_{ji}} = \frac{dE}{do_j} \times \frac{do_j}{dnet_j} \times \frac{dnet_j}{dw_{ji}}$$

$$= \frac{dE}{do_j} \times f'(net_j) \times o_i$$



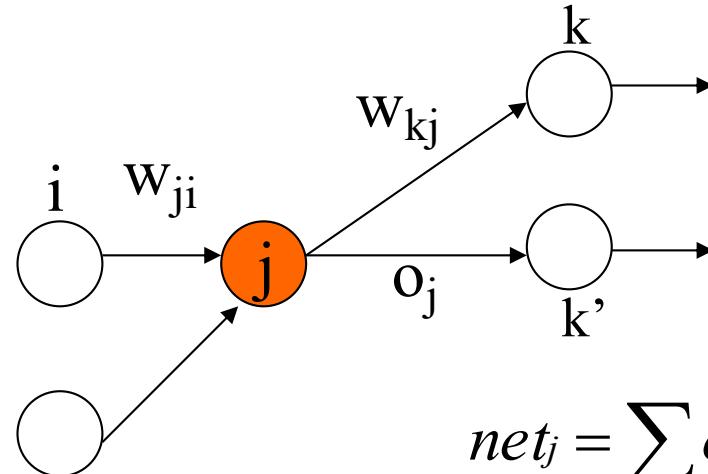
$$net_j = \sum_i o_i w_{ji}$$
$$o_j = f(net_j)$$

Backpropagation – Hidden nodes

If w_{ji} is a **hidden node weight**:

$$\frac{dE}{dw_{ji}} = \frac{dE}{do_j} \times \frac{do_j}{dnet_j} \times \frac{dnet_j}{dw_{ji}}$$

$$= \frac{dE}{do_j} \times f'(net_j) \times o_i$$



$$net_j = \sum_i o_i w_{ji}$$
$$o_j = f(net_j)$$

Note that as j is a hidden node, **we do not know its target**.
Hence, dE/do_j can only be calculated through j 's **contribution to the derivative of E w.r.t net_k at the output nodes**:

$$\frac{dE}{do_j} = \sum_k w_{kj} \times \frac{dE}{dnet_k}$$

dE/dy for Other Loss Functions

Binary Cross Entropy

$$E = - [t \log y + (1-t) \log(1-y)]$$

$$\begin{aligned} dE/dy &= - [t/y - (1-t) \cdot 1/(1-y)] \\ &= -t/y + (1-t)/(1-y) \end{aligned}$$

If $t=1$, derivative is simply $-1/y$.

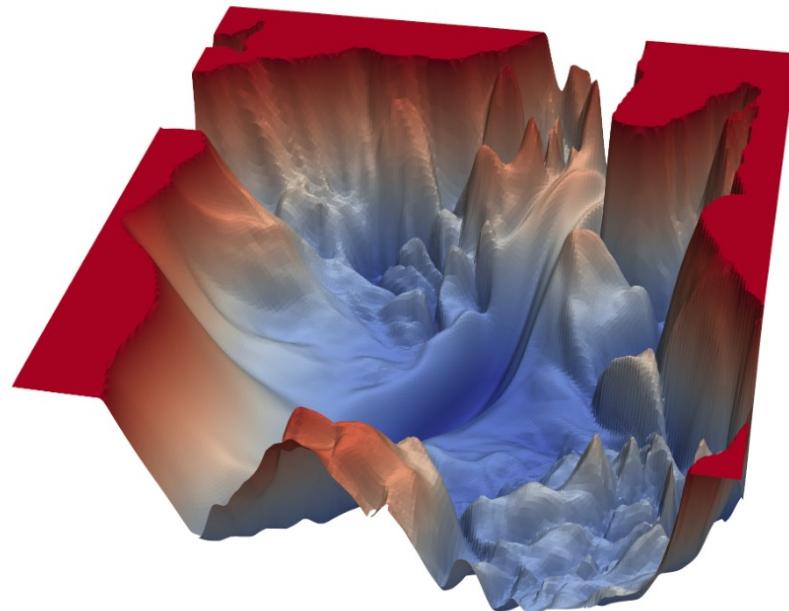
- Since the gradient is negative, increasing y reduces the loss, which is intuitive because we want y to approach 1.

If $t=0$, derivative is simply $1/(1-y)$.

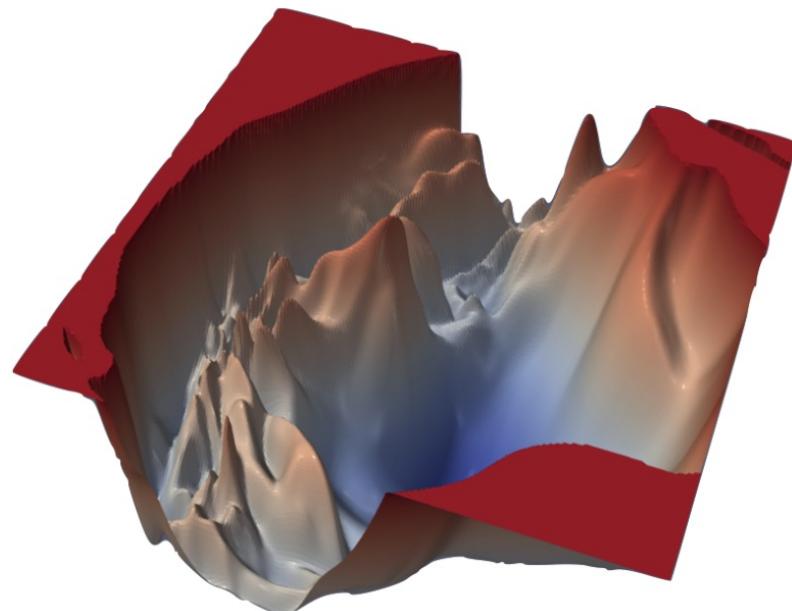
- Since the gradient is positive, decreasing y reduces the loss, which is intuitive because we want y to approach 0.

Error Landscape

VGG-56



VGG-110



Summary

- Gradient descent is the first and standard learning algorithm used in NNs
 - Finds a local minima of the error function
 - Stochastic gradient descent (SGD) can escape local minima
 - Error gradients are always computed using chain rule and propagated backwards – layer by layer
- Be able to compute gradient descent manually