

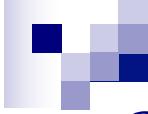
# *Combining Multiple Learners*

## *Part 1*

*Ethem Chp. 15*

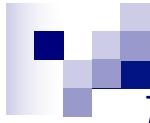
*Hastie Chp 8*

*Haykin Chp. 7, pp. 351-370*



# Overview

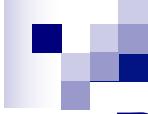
- **Introduction**
  - Motivation
- **Static structures:** the responses of several *experts* (individual networks) are combined in a way that **does not** involve the input signal.
  - *Ensemble averaging / Majority Voting*
  - *Boosting*
  - *Stacking*
  - *Error orrecting Output Codes*
- **Dynamic structures:** the input signal actuates the mechanism that combines the responses of the experts.
  - *Mixture of experts*
  - *Hierarchical mixture of experts*



# Motivation

- When designing a learning machine, we generally make some choices:
  - parameters of machine, training data, representation, etc...
- This implies some sort of **variance** in performance
- **Why not keep all machines and combine their predictions?**
- **Intuition:** Combining experts opinions, votes,...

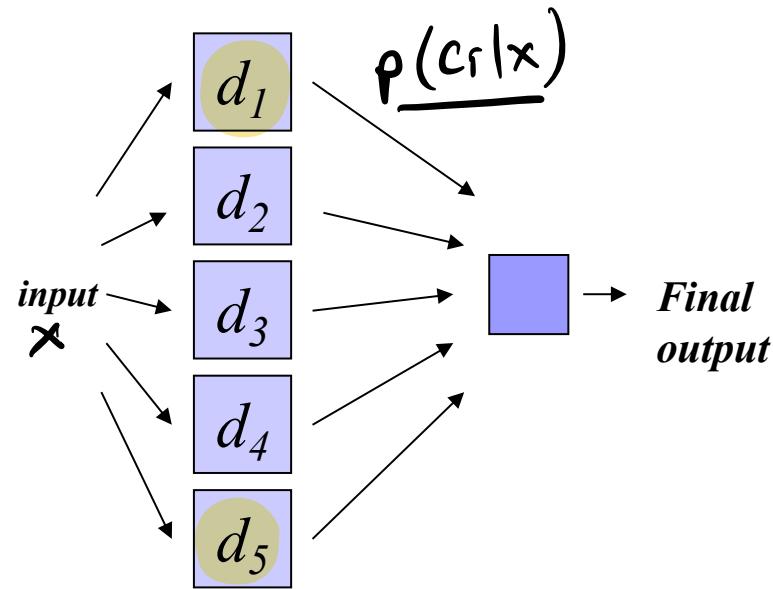
- Ensemble Learning / Classifier Combination:
  - do not learn a single classifier but learn a set of **base classifiers**
  - combine the predictions of multiple classifiers
- Hope to....:
  - **reduce variance:** results are less dependent on peculiarities of a single training set
  - **reduce bias:** a combination of multiple classifiers may learn a more expressive concept class than a single classifier
- Important:
  - formation of an ensemble of **diverse classifiers** from a single training set



## Rationale

- **No Free Lunch thm:** “There is no algorithm that induces the most accurate learner in any domain, all the time.”
  - <http://www.no-free-lunch.org/>
- Generate a group of **base-learners** which when combined has higher accuracy
- Different learners use different
  - **Algorithms:** making different assumptions
  - **Hyperparameters:** e.g number of hidden nodes in NN, k in k-NN
  - **Representations:** diff. features, multiple sources of information
  - **Training sets:** small variations in the sets or diff. subproblems

# Reasons to Combine Learning Machines



**Terminology:**  
Ensemble formed  
from **base classifiers**

Lots of different combination methods:

Most popular are *ensemble averaging* and *majority voting*.

$$\frac{1}{L} \sum_j p(c_{kj}|x)$$

# Ensemble Averaging

- Regression

$$y = \sum_{j=1}^L w_j d_j$$

$$w_j \geq 0 \text{ and } \sum_{j=1}^L w_j = 1$$

$d_j$  is the output of the jth base classifier in regression

$d_{ji}$  is the output of the jth base classifier for ith class in classification

- Classification

$$y_i = \sum_{j=1}^L w_j d_{ji}$$

*inversely*  
 $w_j$  often proportional to error rate of classifier:

Learned over a validation set

# Ensemble Averaging

The output of the ensemble is the (weighted) average of the base classifiers' outputs.

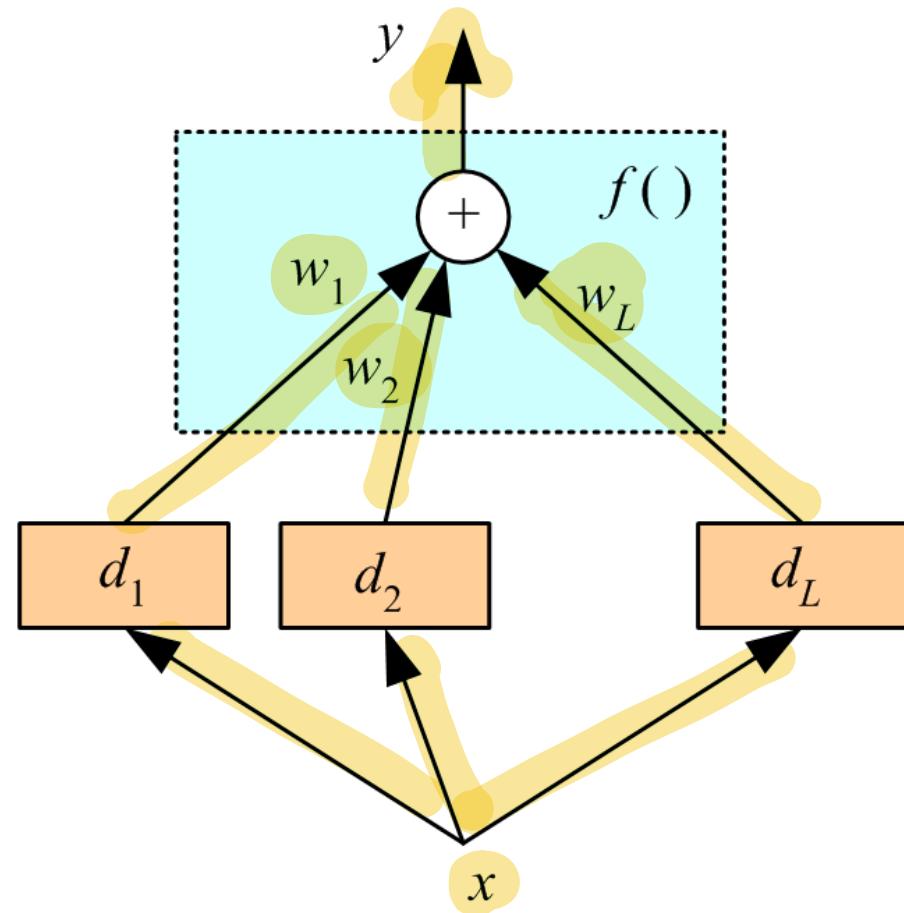
- Regression

$$y = \sum_{j=1}^L w_j d_j$$

$$w_j \geq 0 \text{ and } \sum_{j=1}^L w_j = 1$$

- Classification

$$y_i = \sum_{j=1}^L w_j d_{ji}$$



## Ensemble Averaging

If we use a committee machine  $f_{com}$  whose output is the average:

$$f_{com} = \frac{1}{M} \sum_{i=1}^M d_i$$

*M base learners*

*Error of combination is guaranteed to be lower than the average error:  
ensemble*

$$(f_{com} - t)^2 = \underbrace{\frac{1}{M} \sum_i (d_i - t)^2}_{\text{average error}} - \underbrace{\frac{1}{M} \sum_i (d_i - f_{com})^2}_{\text{always positive}}$$

(Krogh & Vedelsby 1995)

$$\text{Avg}(e_{d_1}, e_{d_2}, e_{d_3})$$

9

- Similarly, we can show that if  $d_j$  are iid:

$$E[f_{com}] = E\left[\sum_j \frac{1}{L} d_j\right] = \frac{1}{L} L \cdot E[d_j] = E[d_j]$$

$$\text{Var}(f_{com}) = \text{Var}\left(\sum_j \frac{1}{L} d_j\right) = \frac{1}{L^2} \text{Var}\left(\sum_j d_j\right) = \frac{1}{L^2} L \cdot \text{Var}(d_j) = \frac{1}{L} \text{Var}(d_j)$$

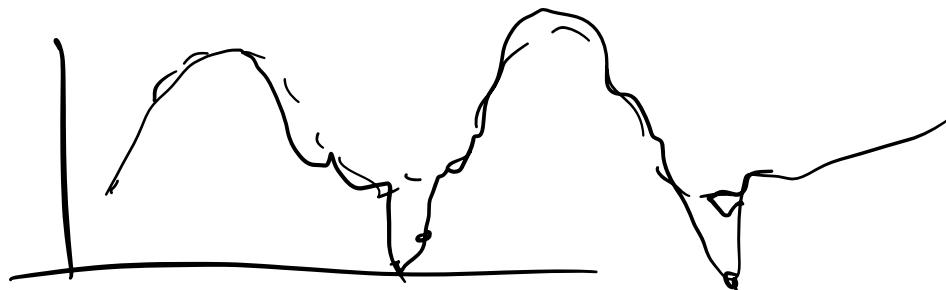
Bias <sup>2</sup> :	$(E_D(d) - f)^2$
Variance:	$E_D[(E_D(d) - d)^2]$

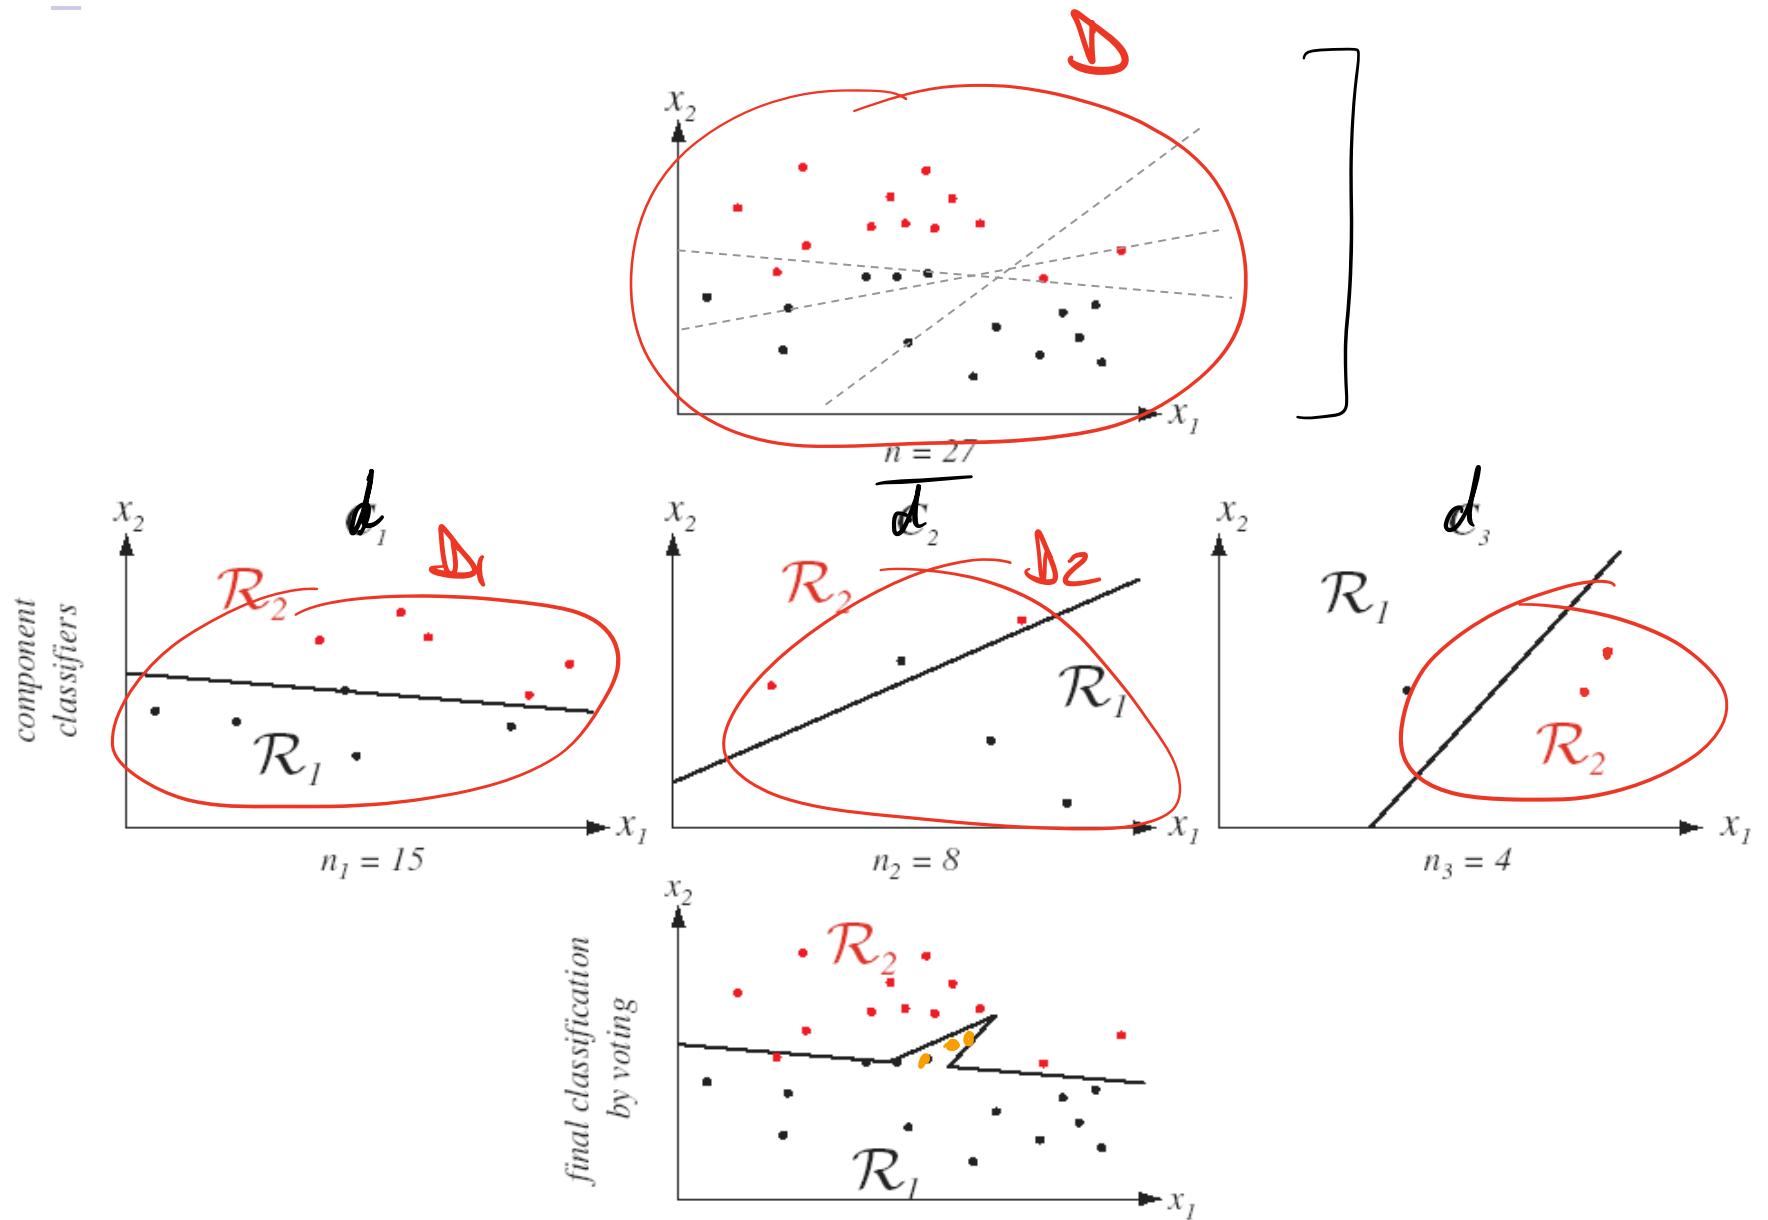
- Bias does not change, variance decreases by  $1/L$

# *Ensemble Averaging*

What we can exploit from this fact:

- Combine multiple experts with the same bias and variance, using ensemble-averaging
  - the bias of the ensemble-averaged system would be **the same as the bias of one of the individual experts**
  - the variance of the ensemble-averaged system would be **less than the variance of one of the individual experts.**
- We can purposefully use complex/flexible models, the variance will be reduced due to averaging.





## Majority Voting

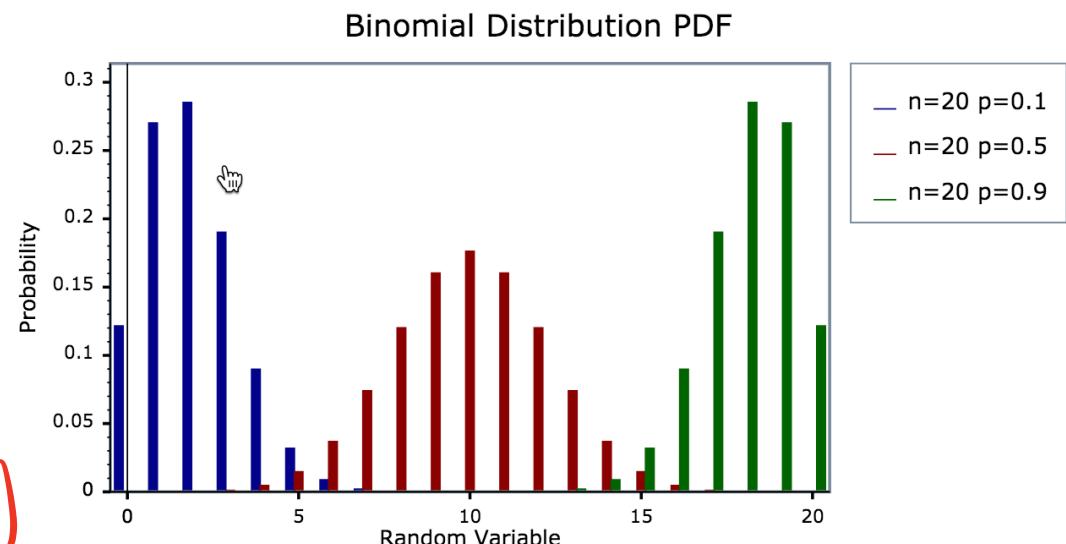
In classification, majority voting chooses the class that gets the most votes from the base classifiers (not averaging the output probabilities).

- Majority voting is also considered a form of ensemble averaging (hard vs soft averaging).

Majority voting ensemble makes an error when more than half of the base classifiers make a mistake. What is this probability?

	<u>Class A</u>	<u>Class B</u>
$d_1$	1	0
$d_2$	1	0
$d_3$	0	1

$\text{ens.} = 1$  (chooses Class A)  
fromm w/ majority vote  
(hard averaging)



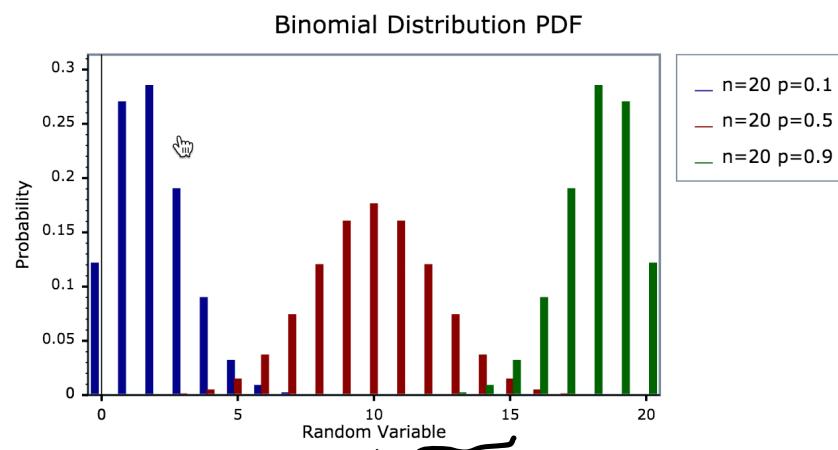
## Reasons to Combine Learning Machines

If the base classifiers are all independent, the probability of error of the **majority voting ensemble** is:

$$P(\text{error}) = \sum_{k=\frac{N}{2}+1}^N \binom{N}{k} p^k (1-p)^{N-k}$$

$p$  : probability  
that  $d_i$   
makes an  
error .

In prev. example,  
sum goes from  
 $k=2$  to  $k=3$ .



## Example

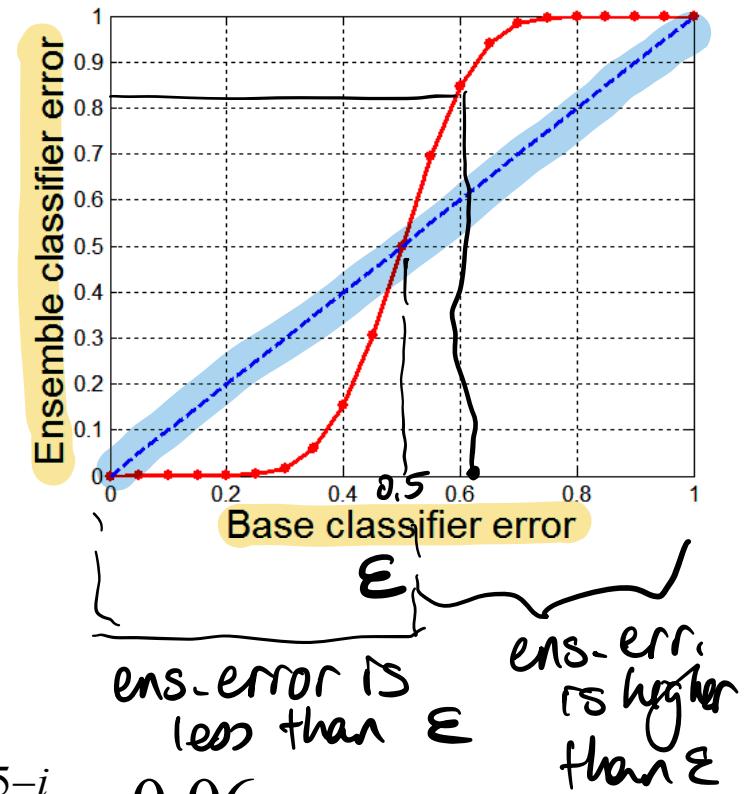
- Suppose there are 25 base classifiers
  - Each classifier has error rate,  $\varepsilon = 0.35$
  - Assume classifiers are independent
    - i.e., probability that a classifier makes a mistake does not depend on whether other classifiers made a mistake
    - Note: in practice they are not independent!
- Probability that the ensemble classifier makes a wrong prediction
  - The ensemble makes a wrong prediction if the majority of the classifiers makes a wrong prediction
  - The probability that 13 or more classifiers err is

$$\sum_{i=13}^{25} \binom{25}{i} \varepsilon^i (1-\varepsilon)^{25-i} \approx 0.06 \ll \varepsilon$$

# Why Majority Voting works?

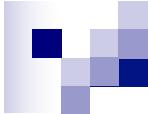
- Suppose there are 25 base classifier, and
  - Each classifier has error rate,  $\varepsilon = 0.35$
  - Errors made by classifiers are uncorrelated
- Probability that the ensemble classifier makes a wrong prediction:

$$P(X \geq 13) = \sum_{i=13}^{25} \binom{25}{i} \varepsilon^i (1-\varepsilon)^{25-i} = 0.06$$



Note: Ensemble does not help when base classifier error is high!  
(where red curve is higher than the blue line)

- We want the base learners to be **complementary**
  - What if they were all the same or very similar?
    - No use...
- **Reasonably accurate**, but not necessarily very accurate



# Overview

- Introduction
  - Rationale
- Combination Methods
  - Static Structures
    - Ensemble averaging / Majority Voting
    - **Bagging**
    - Boosting
    - Error Correcting Output Codes
  - Dynamic structures
    - Mixture of Experts
    - Hierarchical Mixture of Experts

## Ensemble Methods > Bagging

Voting method where base-learners are made different by training over slightly different training sets.

Bagging (Bootstrap Aggregating) - Breiman, 1996

take a training set  $D$ , of size  $N$

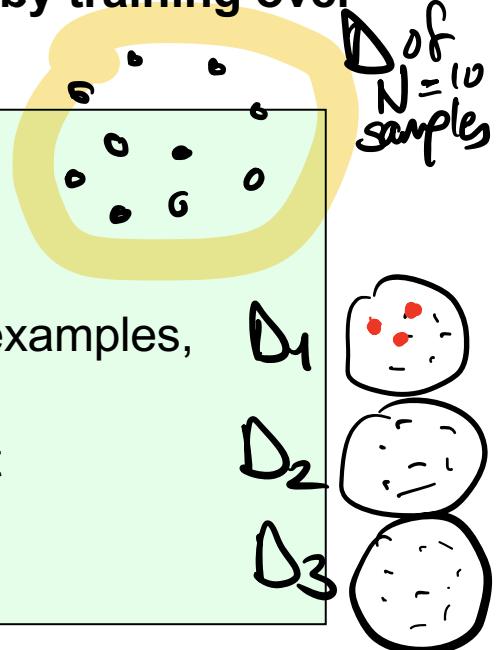
for each network / tree / k-nn / etc... ~~base learner~~

- build a new training set by sampling  $N$  examples, randomly with replacement, from  $D$

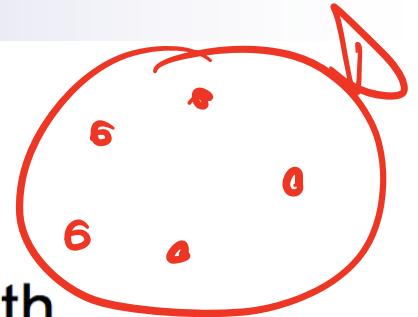
- train your machine with the new dataset

end for

output is average/vote from all machines trained



- Resulting base-learners are similar because they are drawn from the same original sample
- Resulting base-learners are slightly different due to chance



- Generate new training sets using sampling with replacement

→  $D_1$   
 →  $D_2$   
 →  $D_3$

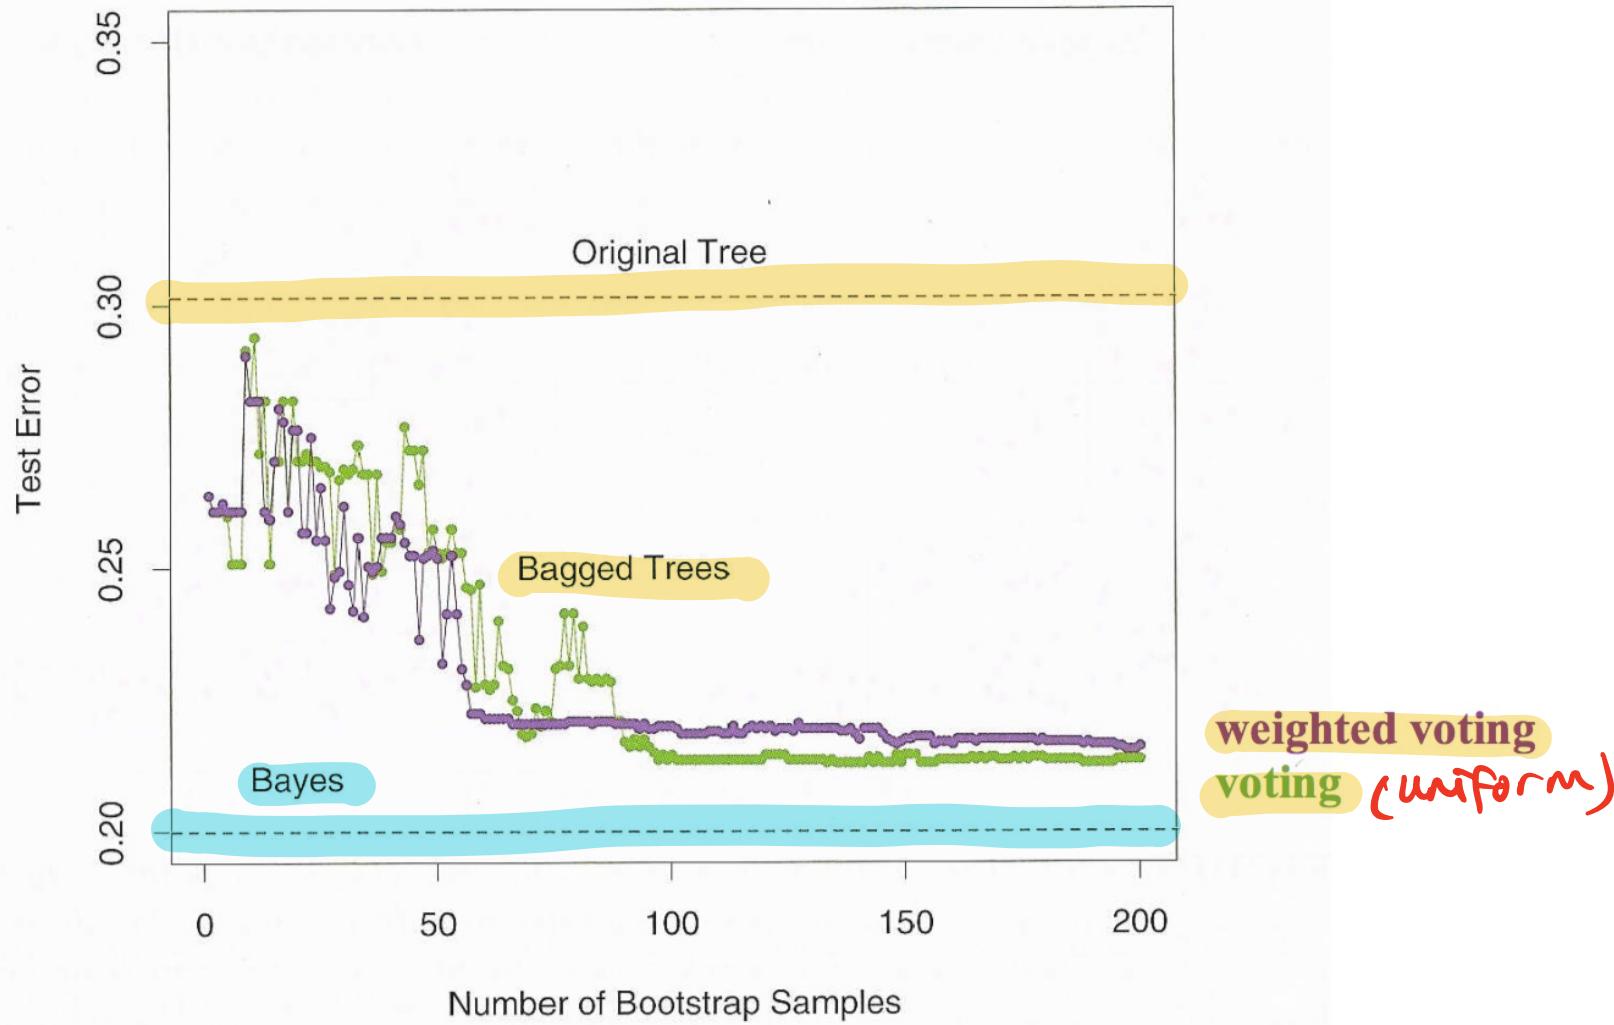
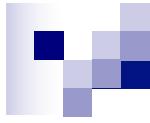
Original Data	1	2	3	4	5	6	7	8	9	10
Bagging (Round 1)	7	8	10	8	2	5	10	10	5	9
Bagging (Round 2)	1	4	9	1	2	3	2	7	3	2
Bagging (Round 3)	1	8	5	10	5	5	9	6	3	7

- some examples may appear in more than one set (samples 1, 3, 4, 6 are not in  $D_1$ )
- for each set, the probability that a given example doesn't appear in it is  $\Pr(x \notin D_i) = (1 - \frac{1}{n})^n$  → 0.3678
- i.e., less than 2/3 of the examples appear in one bootstrap sample

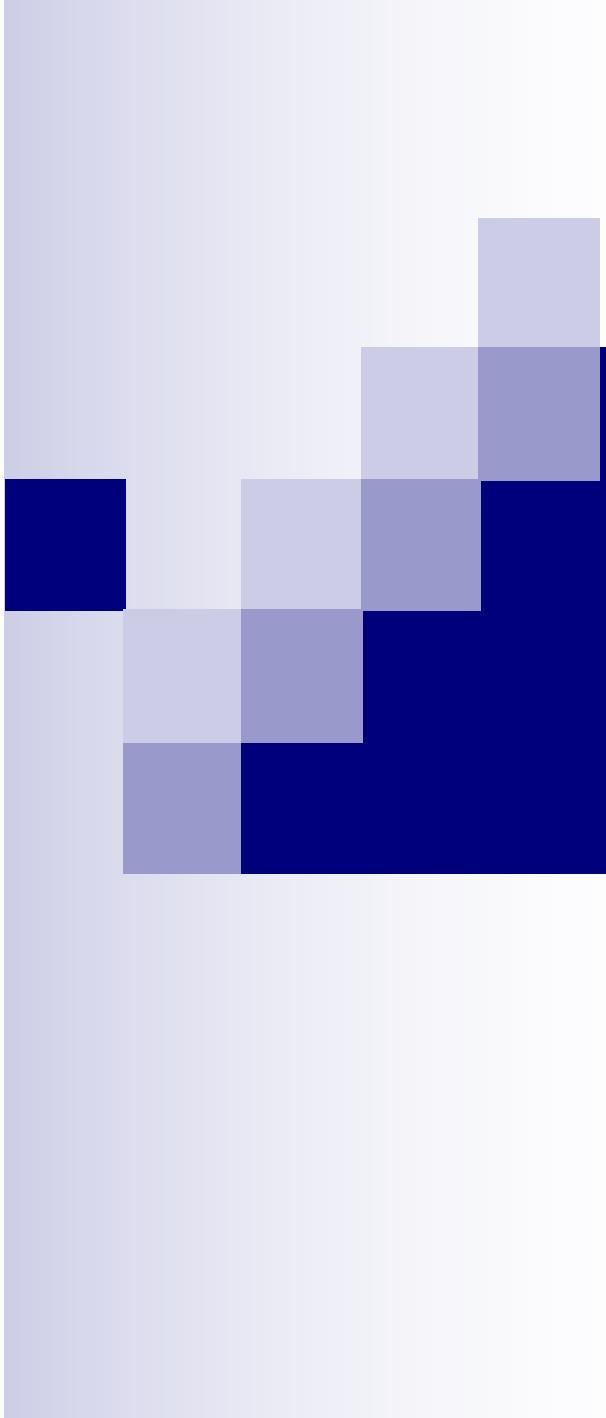


# *Bagging*

- Not all data points will be used for training
  - Waste of training set
    - Each sample has a probability of 0.37 of not being selected in any one bootstrap training set.
    - I.e. only about 2/3rd of the samples are used in any one bootstrap sample.
    - They can be used for testing (see **Out-of-Bag error** in Random Forests)
- Bagging is suitable for **unstable learning** algorithms
  - Unstable algorithms change significantly due to small changes in the data.
  - Such as MLPs, decision trees

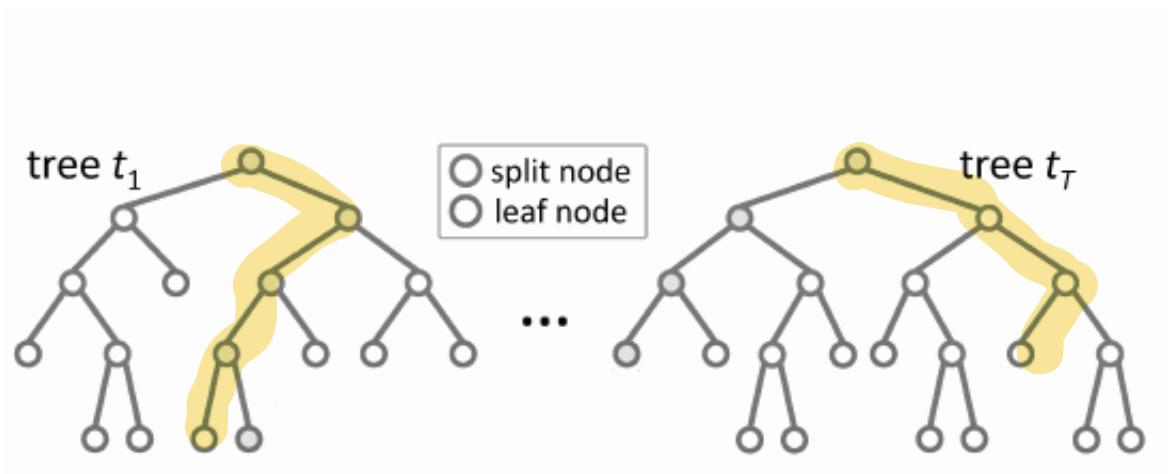


from Hastie, Tibshirani, Friedman: The Elements of Statistical Learning, Springer Verlag 2001

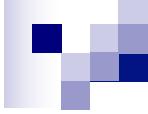


*Random Forests*

## *Random Forests*



- An improved method over **simple bagged trees**
- All trees vote to produce a final answer.



## *Motivation*

- **With decision trees, it is found that optimal cut points can strongly depend on the training set used.**
  - This suggested using multiple trees and use voting to combine the results.
- **Averaging the outputs of the trees reduces overfitting.**
- **For the use of multiple trees to be most effective, the trees should be independent as possible.**
  - Splitting using a random subset of features hopefully achieves this.
  - If the trees really are independent, the performance should improve with more trees

## The Random Forests Algorithm

//Random forest with k trees

Given a training set  $S$

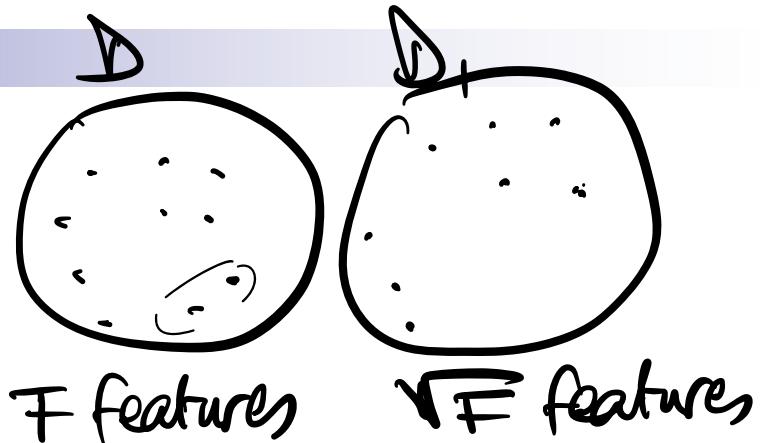
- For  $i = 1$  to  $k$  do:
- Build subset  $S_i$  by sampling with replacement from  $S$   
Learn tree  $T_i$  from  $S_i$  as:

At each node:

Choose best split from random subset of  $F$  features

Each tree grows to the largest extend, no pruning

- Make predictions according to majority vote of the set of  $k$  trees.



- **Bagging**

Generate randomized training sets by sampling with replacement from the full training set (bootstrap sampling)

Full training set      

D <sub>1</sub>	D <sub>2</sub>	D <sub>3</sub>	D <sub>4</sub>	D <sub>5</sub>	D <sub>6</sub>	D <sub>7</sub>	D <sub>8</sub>	D <sub>9</sub>	D <sub>10</sub>	D <sub>11</sub>	D <sub>12</sub>
----------------	----------------	----------------	----------------	----------------	----------------	----------------	----------------	----------------	-----------------	-----------------	-----------------

Random “bag”      

D <sub>4</sub>	D <sub>9</sub>	D <sub>3</sub>	D <sub>4</sub>	D <sub>12</sub>	D <sub>10</sub>	D <sub>10</sub>	D <sub>7</sub>	D <sub>3</sub>	D <sub>1</sub>	D <sub>6</sub>	D <sub>1</sub>
----------------	----------------	----------------	----------------	-----------------	-----------------	-----------------	----------------	----------------	----------------	----------------	----------------

- **Feature subset selection**

Choose different random subsets of the full feature vector to generate each tree

Full feature vector      

f <sub>1</sub>	f <sub>2</sub>	f <sub>3</sub>	f <sub>4</sub>	f <sub>5</sub>	f <sub>6</sub>	f <sub>7</sub>	f <sub>8</sub>	f <sub>9</sub>	f <sub>10</sub>
----------------	----------------	----------------	----------------	----------------	----------------	----------------	----------------	----------------	-----------------

 $F$

Feature subset      

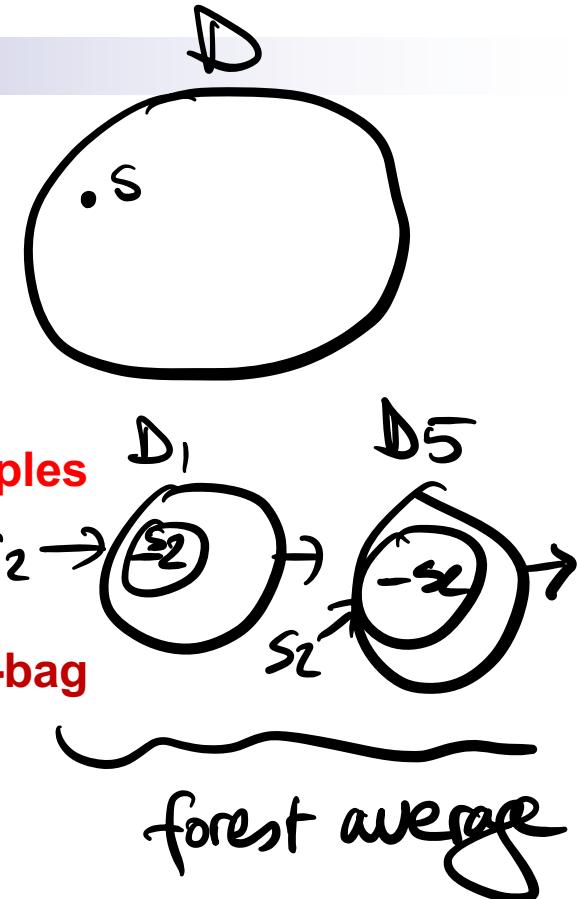
f <sub>4</sub>	f <sub>6</sub>	f <sub>7</sub>	f <sub>9</sub>	f <sub>10</sub>
----------------	----------------	----------------	----------------	-----------------

 $\sqrt{F}$

features

- Typically 5 – 100 trees are used. Often only a few trees are needed.
- Results seem fairly insensitive to the number of random attributes that are tested for each split. A common default is to use the square root of the number of attributes.
- Trees are fast to generate because fewer attributes have to be tested for each split and no pruning is needed.

# Out-of-Bag Error



- To grow one tree
  - Bootstrap sample set from learning set L
  - Remaining samples are called **out-of-bag samples**
- For each sample S of the learning set
  - Look for all the trees for which S was out-of-bag
  - Build the corresponding sub-forest
  - Predict the class of S with it
  - Measure error on S
- **Out-of-bag error** = average over all samples of S
  - Predictions not made using the whole forest... but with some aggregation
- Provides an estimation of the generalization error
  - Can be used to decide when to stop adding trees to the forest

## Features of Random Forests



- It gives some of the best accuracy among current algorithms.
- It is efficient.
- It generates an internal **unbiased estimate of the generalization error** as the forest is built.

Plus...

- It handles **missing data** effectively.
  - A tree doesn't use all the features, plus other methods when a feature value is unknown...
- It gives estimates of **what variables are important** in the classification.
  - Impurity based importance (average info gain from that feature over all trees)
  - Permutation importance
- It has methods for balancing error in class population **unbalanced data sets**.
  - Bootstrap sampling introduces variation in class proportions, or one can use a sample with equal class priors